

Siegiert Relation Violation as a Probe of Spacetime Foam

A. G. Kim*

Physics Division, Lawrence Berkeley National Laboratory, Berkeley, California 94720, USA

P. E. Nugent

*Applied Mathematics and Computational Research Division,
Lawrence Berkeley National Laboratory, Berkeley, California 94720, USA*

L. Wang

Department of Physics and Astronomy, 4242 TAMU, College Station, Texas 77843-4242, USA

(Dated: May 12, 2026)

Quantum gravity predicts that spacetime fluctuates stochastically at the Planck scale — a “foam” that imprints a tiny achromatic arrival-time jitter on light from distant sources. We show that this jitter affects first- and second-order temporal coherence, $g^{(1)}$ and $g^{(2)}$, differently, breaking the Siegert relation $g^{(2)}(\tau) - 1 = |g^{(1)}(\tau)|^2$ that holds for all Gaussian chaotic sources when the foam fluctuates on timescales shorter than the optical coherence time. For spatially extended sources, cross-ray jitter between independent directions additionally suppresses the zero-lag bunching peak $g^{(2)}(0)$ through a mechanism that depends only on $\sigma_\ell^2(D)$ and is independent of the foam’s temporal dynamics. The Siegert violation inherits the wavelength dependence of $g^{(1)}$, while the zero-lag $g^{(2)}(0) < 2$ suppression for extended sources is achromatic and depends on the source angular size in units of the foam transverse coherence scale; combined with their distance scaling, these signatures distinguish foam from plasma and turbulence and are complementary to existing imaging and time-of-flight tests. Together they open a new observational window on the path-length variance predicted by the α -foam model.

I. INTRODUCTION

A long-standing expectation from the union of quantum mechanics and general relativity is that the smooth pseudo-Riemannian manifold of classical spacetime breaks down at the Planck length $\ell_P = \sqrt{\hbar G/c^3} \simeq 1.616 \times 10^{-35}$ m. In Wheeler’s original “spacetime foam” picture [1], quantum fluctuations of the metric on this scale render the geometry stochastic; similar conclusions are reached in the Euclidean path-integral treatment [2] and in modern holographic and entropic-gravity arguments from very different starting points. While a complete theory of quantum gravity remains elusive, the question of whether residual imprints of Planck-scale geometric fluctuations can accumulate over cosmological baselines and become observable is one of the few empirical handles on physics in the Planck regime, and it has motivated a substantial body of phenomenological work over the past two decades [3, 4].

A widely studied parametrization, due to Ng and collaborators [4–6] and referred to here as the α -foam model, assigns the cumulative line-of-sight path-length variance acquired by light traversing a distance D the power-law form

$$\sigma_\ell^2(D) = A_\alpha D^{2(1-\alpha)} \ell_P^{2\alpha}, \quad (1)$$

where A_α is a dimensionless coefficient of order unity and the exponent α encodes the underlying foam dynamics. Three values are conventionally distinguished. The “conservative” case $\alpha = 1$ follows from Wheeler’s original dimensional argument and gives $\sigma_\ell \sim \ell_P$ regardless of D – observationally inconsequential. The “holographic” case $\alpha = 2/3$ is motivated by the requirement that the information content of a region of spacetime saturate the Bekenstein–Hawking bound and scale with the area of its boundary in Planck units [5]; it predicts $\sigma_\ell \sim (D\ell_P^2)^{1/3}$, of order 10^{-15} m over 1 Gpc. The “random-walk” case $\alpha = 1/2$ treats successive Planck cells as independent path-length kicks and gives $\sigma_\ell \sim (D\ell_P)^{1/2}$, roughly 10^{-5} m over the same distance – large enough to randomize the optical phase of light from sources at $z \gtrsim 1$. The three cases bracket the range investigated most thoroughly; for a recent review situating them within the broader quantum-gravity phenomenology program, see Carlip [3].

Two complementary observational channels have been used to test these models. The first exploits the fact that random Planck-scale path-length deviations imprint a stochastic phase $\sigma_\phi = 2\pi\sigma_\ell/\lambda$ across the aperture of any telescope, reducing the Strehl ratio by $S \approx \exp(-\sigma_\phi^2)$ and degrading the diffraction-limited point-spread function.

* agkim@lbl.gov

Detections of crisp, near-diffraction-limited images of distant active galactic nuclei with the Hubble Space Telescope and ground-based instruments [7–10], together with stellar interferometry, exclude $\alpha \lesssim 2/3$ at optical wavelengths for $D \sim 1\text{--}10$ Gpc, leaving the holographic and conservative cases as the only quantitatively viable scalings. These exclusions assume transverse foam correlations are negligible on aperture scales; models with long-range transverse correlations can evade blurring constraints while retaining large $\sigma_\ell(D)$ [11]. The second channel uses energy-dependent time-of-flight measurements of high-energy astrophysical transients: gamma-ray bursts observed by Fermi [12] and TeV blazar flares observed by H.E.S.S. [13] constrain the linear and quadratic dispersion that some quantum-gravity models predict, with LHAASO observations of GRB 221009A [14] providing the current leading limits. Together the two channels have winnowed the parameter space but not eliminated it, and several proposed scalings – including the holographic $\alpha = 2/3$ – remain unconstrained at any wavelength accessible to current instrumentation. A recent analysis of GRB 221009A claims consistency with holographic foam blurring [15], though without formal statistical significance.

Both existing channels share a structural limitation that motivates the present work. The imaging tests are essentially binary: a wavefront either remains coherent enough to image or it does not, so that once a positive detection is registered the constraint saturates and no further information is recovered about the magnitude of the residual phase variance. The time-of-flight tests, conversely, rely on a chromatic dispersion signature that is precisely what is *not* expected from a geometric-optics path-length fluctuation, for which the timing offset $\delta t = \Delta\ell/c$ is frequency-independent. The Fermi, H.E.S.S., and LHAASO analyses specifically seek a Lorentz-invariance-violating (LIV) dispersion of the form $\delta t \propto E^n/E_{\text{QG}}^n$: a deterministic, energy-ordered delay. A stochastic, energy-independent geometric jitter produces no such signal, so the absence of a detected LIV dispersion in GRB or blazar data leaves the foam entirely unconstrained by those measurements. A foam model whose accumulated phase variance lies below the imaging threshold but whose timing jitter is achromatic is therefore invisible to both channels. Whether observables exist that bridge this gap – sensitive to a phase variance too small to erase fringes outright, but achromatic and thus distinguishable from plasma scintillation or instrumental dispersion – is a natural question, and it is the question we address here.

We develop a framework based on the first- and second-order coherence functions $g^{(1)}(\tau)$ and $g^{(2)}(\tau)$ of light from astrophysical sources. The first-order coherence – the field–field autocorrelation, equivalently the fringe visibility at delay τ – inherits a foam suppression that interpolates continuously between unity and the Lieu–Hillman fringe-erasure limit [7], recovering the Strehl and fringe-visibility tests as limiting cases of the same statistic. The second-order coherence – the intensity–intensity autocorrelation measured in a Hanbury Brown–Twiss experiment [16] – is to leading order blind to the carrier-phase modulation that suppresses $g^{(1)}$, but acquires a distinct signature from the differential timing jitter between photon arrivals. The two together violate the Siegert relation $g^{(2)}(\tau) - 1 = |g^{(1)}(\tau)|^2$ [17] in a foam-specific way, providing a discriminant that does not require the intrinsic source visibility to be modeled. The Siegert violation requires the foam to fluctuate on timescales shorter than the optical coherence time; the zero-lag bunching suppression $\delta g^{(2)}$ for extended sources, by contrast, depends only on $\sigma_\ell^2(D)$ and is independent of the foam’s temporal dynamics. For spatially resolved sources, transverse foam coherence produces an additional, geometrically distinct effect: a suppression of the zero-lag HBT bunching peak that depends on the source angular size and provides a direct probe of the transverse foam coherence length $r_\perp$ – a degree of freedom accessible via single-aperture second-order temporal coherence. Next-generation sensitivity is achievable via two complementary paths: large-aperture single telescopes measuring the HBT effect, and multi-telescope intensity interferometer arrays such as VERITAS [18, 19] – both increase collecting area and improve signal-to-noise for nearby resolved sources, while the temporal $g^{(1)}$ channel is accessible to high-resolution spectroscopy of distant quasars. Detection of both point-like and extended-source signals at the same redshift would jointly probe $\{\alpha, \sigma_\ell(D), r_\perp\}$: the point-source Siegert discrepancy is sensitive to α and D with no sensitivity to $r_\perp$, while the zero-lag bunching suppression for an extended source is sensitive to $r_\perp$ largely independently of the intrinsic source spectrum. Crucially, their consistency provides a test of whether any observed signatures are jointly consistent with the predictions of the foam model, rather than a common instrumental or astrophysical systematic.

The remainder of the paper is organized as follows. Section II develops the model: the spatiotemporal covariance of the foam-induced path-length fluctuation field, the differential-jitter distribution at a single ray, and its generalization to ray pairs. Section III computes the foam suppression of first-order coherence, derives closed-form results for Lorentzian and Voigt line profiles, and connects it to the published Strehl and fringe-visibility tests. Section IV develops the second-order coherence, including the conditional Siegert relation and its narrow-band limit. Section V quantifies the resulting Siegert violation. Section VI extends the analysis to spatially resolved sources, derives the suppression of the zero-lag HBT bunching peak evaluated for Lorentzian and Voigt profiles, identifies the distance scaling and achromaticity as discriminants against conventional propagation phase screens, and shows how detection of both point-like and extended-source signals would jointly probe $\{\alpha, \sigma_\ell(D), r_\perp\}$.

II. MODEL

A. Arrival-Time Jitter as a Time Reparameterization

We model the detected field as a stochastically time-warped version of the intrinsic source field:

$$E_{\text{obs}}(t) = E_{\text{src}}(t - \delta t(t)), \quad (2)$$

where $E_{\text{src}}(t)$ is a single spatial mode of a stationary zero-mean complex Gaussian process (chaotic light, appropriate for virtually all known astrophysical sources) and $\delta t(t)$ is the foam-induced arrival-time jitter, a stationary stochastic process statistically independent of E_{src} . The jitter arises from cumulative path-length fluctuations along the line of sight of propagation distance D :

$$\delta t(t) = \frac{1}{c} \Delta \ell_{\text{tot}}(t). \quad (3)$$

In the geometric-optics limit the path-length shift $\Delta \ell_{\text{tot}}$ is a purely geometric quantity and hence frequency-independent (achromatic), in contrast to plasma dispersion which scales as λ^2 . This achromaticity is a key discriminant between foam and conventional propagation effects.

B. Spatial and Temporal Fluctuations of Spacetime Foam

Following Ng and collaborators [4–6] (see also Ref. [3] for a review of foam models and their observational constraints), the single-ray path-length variance obeys the power-law scaling

$$\sigma_\ell^2(D) = A_\alpha D^{2(1-\alpha)} \ell_P^{2\alpha}, \quad (4)$$

where α parametrizes the foam model ($\alpha = 1$: conservative; $\alpha = 2/3$: holographic; $\alpha = 1/2$: random walk), and A_α is a dimensionless coefficient of order unity.

Because foam is a property of spacetime rather than of any preferred direction, the same dimensional logic that governs longitudinal path-length fluctuations motivates the existence of a characteristic transverse coherence length $r_\perp$. However, the longitudinal and transverse cases are geometrically distinct: in the longitudinal case D counts the number of cells traversed along the ray, while in the transverse case b is a separation between rays that share cells along the line of sight — the number of shared cells depends on both b and D in a way that is not captured by a simple dimensional analogy. Accordingly, $r_\perp$ is treated as an independent free parameter in Section VI. We denote the normalized transverse coherence function by $\Phi(b/r_\perp)$, satisfying $\Phi(0) = 1$ and $\Phi(b/r_\perp) \rightarrow 0$ for $b \gg r_\perp$. The functional form of Φ depends on the specific foam model. For the pixellon model, Lee, Zurek & Chen [11] find that the transverse correlation structure involves long-range angular correlations (a Legendre polynomial series in aperture-plane angle) that do not reduce to a simple $\Phi(b/r_\perp)$ ansatz, but whose long-range character strongly suppresses image blurring relative to the uncorrelated case. In the present framework Φ is treated as a free function; a foam signal would manifest as an angular-size dependence of $\delta g^{(2)}$ (Section VI).

Turning to the temporal evolution of spacetime foam — an aspect not specified by Eq. (4) — quantum foam models posit that each Planck-scale cell along the line of sight fluctuates in time. Two photons traversing the same cell at coordinate times separated by τ experience path-length fluctuations that are correlated for $\tau \ll \tau_{\text{foam}}$ but decorrelated for $\tau \gg \tau_{\text{foam}}$, where τ_{foam} is the characteristic temporal fluctuation timescale at each cell. We introduce $C_\alpha(\tau; D)$ as the normalized *temporal* autocorrelation of the per-cell path-length fluctuation, integrated over the $\sim D/\ell_P$ cells along the line of sight (where $\ell_P = \sqrt{\hbar G/c^3} \simeq 1.616 \times 10^{-35}$ m is the Planck length). The timescale τ_{foam} is an additional free parameter of the foam model, not fixed by the α -model variance relation alone. If τ_{foam} is of order the Planck time, $\ell_P/c \approx 5 \times 10^{-44}$ s, then $C_\alpha(\tau) \approx 0$ for all accessible lags: the foam is fully decorrelated on any observable timescale, the most optimistic scenario for detection. In the opposite *quasi-static* limit $\tau_{\text{foam}} \gg D/c$ (foam coherent across the entire line of sight), $C_\alpha(\tau) \approx 1$ for all τ of interest and every temporal observable vanishes. The parameters A_α , $C_\alpha(\tau; D)$, τ_{foam} , $r_\perp$, and Φ_{eff} are convenient parameterizations of distinct physical effects — single-ray variance, temporal foam dynamics, transverse coherence, and source angular structure — and each governs a distinct aspect of the signal's detectability.

The full spatiotemporal covariance between two rays separated by transverse distance b takes the factorized form

$$\langle \Delta \ell_{\text{tot}}(t; \boldsymbol{\theta}) \Delta \ell_{\text{tot}}(t + \tau; \boldsymbol{\theta}') \rangle = \sigma_\ell^2(D) C_\alpha(\tau; D) \Phi\left(\frac{b(\boldsymbol{\theta}, \boldsymbol{\theta}')}{r_\perp}\right), \quad (5)$$

where $b(\boldsymbol{\theta}, \boldsymbol{\theta}') \simeq D|\boldsymbol{\theta} - \boldsymbol{\theta}'|$. The single-ray temporal covariance appropriate for a point source is the special case $\boldsymbol{\theta} = \boldsymbol{\theta}'$ (so $b = 0$):

$$\langle \Delta \ell_{\text{tot}}(t) \Delta \ell_{\text{tot}}(t + \tau) \rangle = \sigma_\ell^2(D) C_\alpha(\tau; D). \quad (6)$$

This product form — separating the total variance from the normalized spatiotemporal shape — is a structural assumption, not a consequence of Eq. (4), that defines a tractable class of foam models within which the analytic framework is exact. Observationally, factorization predicts that the ratio $[g_{\text{obs}}^{(2)}(\tau) - 1]_{\text{extended}}/[g_{\text{obs}}^{(2)}(\tau) - 1]_{\text{point}}$ equals Φ_{eff} and is independent of both τ and D ; any dependence of this ratio on lag or source redshift would falsify the factorization ansatz. The two principal observables developed below — the Siegert discrepancy $\mathcal{S}(\tau)$ and the zero-lag bunching suppression $\delta g^{(2)}$ — both descend from the single scaling relation Eq. (4), with Sections III–V probing its longitudinal projection and Section VI its transverse projection.

C. Differential Jitter and Its Distribution

With the spatiotemporal covariance model in place, we define the key derived quantity that enters the coherence integrals. Define the *cross-ray differential jitter* between a photon arriving along direction $\boldsymbol{\theta}$ at time $t + \tau$ and a photon arriving along $\boldsymbol{\theta}'$ at time t :

$$\Delta_{\boldsymbol{\theta}, \boldsymbol{\theta}'}(\tau) \equiv \delta t(t + \tau; \boldsymbol{\theta}) - \delta t(t; \boldsymbol{\theta}'). \quad (7)$$

For fixed D , stationarity implies the statistics depend only on τ and the transverse separation $b(\boldsymbol{\theta}, \boldsymbol{\theta}') = D|\boldsymbol{\theta} - \boldsymbol{\theta}'|$. Expanding the square and applying Eq. (5):

$$\begin{aligned} \sigma_\Delta^2(\tau; \boldsymbol{\theta}, \boldsymbol{\theta}') &= \frac{1}{c^2} \langle [\Delta \ell_{\text{tot}}(t + \tau; \boldsymbol{\theta}) - \Delta \ell_{\text{tot}}(t; \boldsymbol{\theta}')]^2 \rangle \\ &= \frac{1}{c^2} \left[\sigma_\ell^2 + \sigma_\ell^2 - 2\sigma_\ell^2 C_\alpha(\tau) \Phi\left(\frac{b}{r_\perp}\right) \right] \\ &= \frac{2\sigma_\ell^2(D)}{c^2} \left[1 - C_\alpha(\tau; D) \Phi\left(\frac{b(\boldsymbol{\theta}, \boldsymbol{\theta}')}{r_\perp}\right) \right]. \end{aligned} \quad (8)$$

For $b \gg r_\perp$, $\Phi \rightarrow 0$ and the variance becomes $2\sigma_\ell^2/c^2$ at all lags, including $\tau = 0$ — a new effect absent in the point-source case, and one that is independent of the temporal fluctuation structure of the foam (i.e. independent of C_α). In the dynamical limit $C_\alpha \rightarrow 0$ this result holds for all separations b , independently of Φ .

By the central limit theorem (CLT), $\Delta_{\boldsymbol{\theta}, \boldsymbol{\theta}'}(\tau)$ is a sum of contributions from many foam cells and is therefore Gaussian with zero mean:

$$\Delta_{\boldsymbol{\theta}, \boldsymbol{\theta}'}(\tau) \sim \mathcal{N}(0, \sigma_\Delta^2(\tau; \boldsymbol{\theta}, \boldsymbol{\theta}')), \quad (9)$$

where σ_Δ^2 is taken from Eq. (8).

The Gaussian approximation requires the elementary foam cells to be approximately independent. For the cross-ray case the two rays share foam cells along the line of sight, and independence requires $b \gg r_\perp$. In the intermediate regime $b \sim r_\perp$ (equivalently $r_\perp \sim D\Delta\theta$), the cross-ray Gaussian approximation and the Φ_{eff} approximation of Section VID break down together; both introduce errors of the same order and must be assessed jointly.

III. FIRST-ORDER COHERENCE

In Sections III–V we restrict to a point source ($b = 0$), for which the cross-ray differential jitter of Eq. (7) reduces to the single-ray differential jitter

$$\Delta(\tau) \equiv \delta t(t + \tau) - \delta t(t), \quad (10)$$

with variance obtained from Eq. (8) at $\Phi(0) = 1$:

$$\sigma_\Delta^2(\tau; D) = \frac{2\sigma_\ell^2(D)}{c^2} [1 - C_\alpha(\tau; D)]. \quad (11)$$

A. Jitter Convolution Formula for $g^{(1)}$

The observed first-order coherence is

$$g_{\text{obs}}^{(1)}(\tau) = \frac{\langle E_{\text{obs}}^*(t) E_{\text{obs}}(t + \tau) \rangle}{\langle |E_{\text{obs}}(t)|^2 \rangle}. \quad (12)$$

Substituting Eq. (2) and letting $u = t - \delta t(t)$, $v = t + \tau - \delta t(t + \tau)$, so that $v - u = \tau - \Delta(\tau)$, stationarity of E_{src} gives

$$\langle E_{\text{src}}^*(u) E_{\text{src}}(v) \rangle = \langle |E_{\text{src}}|^2 \rangle g_{\text{src}}^{(1)}(\tau - \Delta(\tau)). \quad (13)$$

Since δt is independent of E_{src} , averaging over the jitter yields

$$g_{\text{obs}}^{(1)}(\tau) = \langle g_{\text{src}}^{(1)}(\tau - \Delta(\tau)) \rangle_{\delta t} = \int_{-\infty}^{\infty} \mathcal{N}(\Delta; 0, \sigma_{\Delta}^2(\tau)) g_{\text{src}}^{(1)}(\tau - \Delta) d\Delta. \quad (14)$$

The observed first-order coherence is the intrinsic coherence function convolved with the Gaussian differential-jitter kernel.

B. Closed Forms for Lorentzian and Voigt Sources

As a representative case admitting a closed-form result, we take a Lorentzian line profile, which corresponds to an exponential coherence function via the Wiener–Khinchin theorem:

$$g_{\text{src}}^{(1)}(\tau) = e^{-|\tau|/\tau_c}, \quad (15)$$

where τ_c is the coherence time. Evaluating Eq. (14) by completing the square and splitting on the sign of $\tau - \Delta$ gives, for $\tau > 0$:

$$g_{\text{obs}}^{(1)}(\tau) = e^{\sigma_{\Delta}^2/2\tau_c^2} \cdot \frac{1}{2} \left[e^{-\tau/\tau_c} \operatorname{erfc}\left(\frac{\sigma_{\Delta}/\tau_c - \tau/\sigma_{\Delta}}{\sqrt{2}}\right) + e^{\tau/\tau_c} \operatorname{erfc}\left(\frac{\sigma_{\Delta}/\tau_c + \tau/\sigma_{\Delta}}{\sqrt{2}}\right) \right]. \quad (16)$$

In the weak-jitter limit $\sigma_{\Delta} \ll \tau_c$, the Gaussian average of $e^{-|\tau - \Delta|/\tau_c}$ over $\Delta \sim \mathcal{N}(0, \sigma_{\Delta}^2)$ is dominated by $|\Delta| \ll |\tau|$, so one expands around $\Delta = 0$:

$$g_{\text{obs}}^{(1)}(\tau) \approx e^{-|\tau|/\tau_c} \left[1 + \frac{\sigma_{\Delta}^2(\tau)}{2\tau_c^2} + \mathcal{O}(\sigma_{\Delta}^4) \right]. \quad (17)$$

Here $\sigma_{\Delta}(\tau)$ is τ -dependent through $C_{\alpha}(\tau)$, and the expansion is in $\sigma_{\Delta}(\tau)/\tau_c$ at fixed τ . The correction is positive because $e^{-|\tau|/\tau_c}$ is convex and Jensen's inequality applies. The normalization $g_{\text{obs}}^{(1)}(0) = 1$ is preserved exactly.

For a Voigt profile $g_{\text{src}}^{(1)}(\tau) = e^{-|\tau|/\tau_L} e^{-\tau^2/2\tau_G^2}$ — a Lorentzian of width τ_L convolved with a Gaussian of width τ_G modelling Doppler and turbulent broadening — the jitter kernel combines with the source Gaussian into an effective width and mean,

$$\tilde{\sigma}_1^2 = \frac{\sigma_{\Delta}^2 \tau_G^2}{\sigma_{\Delta}^2 + \tau_G^2}, \quad \mu_1(\tau) = \frac{\tau_G^2 \tau}{\sigma_{\Delta}^2 + \tau_G^2}, \quad (18)$$

giving, for $\tau > 0$:

$$g_{\text{obs}}^{(1)}(\tau) = \frac{\tilde{\sigma}_1}{\sigma_{\Delta}} e^{-\tau^2/2(\sigma_{\Delta}^2 + \tau_G^2)} e^{\tilde{\sigma}_1^2/2\tau_L^2} \frac{1}{2} \left[e^{-\mu_1/\tau_L} \operatorname{erfc}\left(\frac{\tilde{\sigma}_1/\tau_L - \mu_1/\tilde{\sigma}_1}{\sqrt{2}}\right) + e^{\mu_1/\tau_L} \operatorname{erfc}\left(\frac{\tilde{\sigma}_1/\tau_L + \mu_1/\tilde{\sigma}_1}{\sqrt{2}}\right) \right]. \quad (19)$$

The normalization $g_{\text{obs}}^{(1)}(0) = 1$ is preserved: as $\sigma_{\Delta} \rightarrow 0$, $\tilde{\sigma}_1/\sigma_{\Delta} \rightarrow 1$, the exponential prefactors approach unity, and the erfc combination gives 1. In the limit $\tau_G \rightarrow \infty$ (pure Lorentzian), $\tilde{\sigma}_1 \rightarrow \sigma_{\Delta}$ and $\mu_1 \rightarrow \tau$, so Eq. (19) reduces to Eq. (16). In the limit $\tau_L \rightarrow \infty$ (pure Gaussian), the erfc terms combine to unity and $g_{\text{obs}}^{(1)}(\tau) = (\tilde{\sigma}_1/\sigma_{\Delta}) e^{-\tau^2/2(\sigma_{\Delta}^2 + \tau_G^2)}$: the observed coherence is a Gaussian whose width is broadened in quadrature by the jitter.

C. Narrow-Band Limit

For a quasi-monochromatic source with center frequency ω_0 and angular-frequency bandwidth $\Delta\omega \ll \omega_0$, write $E_{\text{src}}(t) = A(t) e^{-i\omega_0 t}$ where $A(t)$ varies on the coherence timescale $\tau_c = 1/\Delta\omega = \lambda^2/(2\pi c \Delta\lambda)$. When $\sigma_\ell/(c\tau_c) \ll 1$, typical realizations of δt leave the envelope unaffected, $A(t - \delta t) \approx A(t)$, and the jitter acts purely as a carrier phase modulation; the observed field takes the form

$$E_{\text{obs}}(t) \approx A(t) e^{-i\omega_0 t} e^{i\omega_0 \delta t(t)}. \quad (20)$$

(Corrections to the envelope approximation are of order $(\sigma_\ell/c\tau_c)^2$; for $\alpha = 1/2$ at $D = 1$ Gpc, $\lambda = 500$ nm, $\Delta\lambda/\lambda = 10^{-3}$ this parameter is ≈ 0.28 , so the exact convolution Eq. (16) should be used instead. For $\alpha = 2/3$ the analogous parameter is $\sim 10^{-11}$ (Table I), so Eq. (22) is exact to any practical precision in the holographic case.) Since $A(t)$ and $\delta t(t)$ are independent, the first-order coherence factorizes. The jitter contribution is the characteristic function of $\Delta(\tau)$ evaluated at ω_0 ; since $\Delta(\tau)$ is Gaussian (Section II C),

$$\langle e^{i\omega_0 \Delta(\tau)} \rangle_{\delta t} = e^{-\omega_0^2 \sigma_\Delta^2(\tau)/2}. \quad (21)$$

This holds when $\tau \gg \tau_{\text{foam}}$ so that foam-cell contributions are independent; in the dynamical limit $\tau_{\text{foam}} \ll \tau_c$ the condition is satisfied at all observationally relevant lags. Substituting $\omega_0 = 2\pi c/\lambda$ and Eq. (11):

$$g_{\text{obs}}^{(1)}(\tau) = g_{\text{src}}^{(1)}(\tau) \cdot \exp(-\sigma_\phi^2 [1 - C_\alpha(\tau; D)]), \quad (22)$$

where we define the rms accumulated phase

$$\sigma_\phi \equiv \frac{2\pi \sigma_\ell(D)}{\lambda} = \frac{2\pi \sqrt{A_\alpha} D^{1-\alpha} \ell_P^\alpha}{\lambda}. \quad (23)$$

We write the foam suppression factor as

$$\mathcal{F}_{\text{foam}}(\tau) \equiv e^{-\sigma_\phi^2 [1 - C_\alpha(\tau; D)]}, \quad (24)$$

so $g_{\text{obs}}^{(1)}(\tau) = g_{\text{src}}^{(1)}(\tau) \cdot \mathcal{F}_{\text{foam}}(\tau)$. At $\tau = 0$: $\mathcal{F}_{\text{foam}}(0) = 1$. At $\tau \rightarrow \infty$: $\mathcal{F}_{\text{foam}}(\infty) = e^{-\sigma_\phi^2 [1 - C_\alpha(\infty; D)]}$, which equals $e^{-\sigma_\phi^2}$ in the dynamical limit ($C_\alpha \rightarrow 0$) and approaches unity for quasi-static foam ($C_\alpha \approx 1$).

D. Regime Summary

Eq. (14) is always valid (exact within the model assumptions); Eq. (16) is its closed form for a Lorentzian source and Eq. (19) for a Voigt source, both likewise exact. Eq. (17) (weak-jitter expansion) requires $\sigma_\Delta(\tau)/\tau_c \ll 1$ at each lag of interest. Eq. (22) (narrow-band, carrier-phase result) requires $\sigma_\ell/(c\tau_c) \ll 1$; it captures the dominant physics when $\sigma_\phi \gtrsim 1$ but $\sigma_\phi \ll \lambda/\Delta\lambda$. These conditions are compatible: the narrow-band condition controls whether the amplitude envelope is displaced by the jitter, while the weak-jitter condition controls whether the carrier-phase suppression is small. When neither approximation applies (i.e. $\sigma_\ell/(c\tau_c) \sim 1$), the full convolution Eq. (14) must be evaluated.

E. Connection to Fringe Visibility and Strehl Ratio

The fringe visibility and Strehl ratio are spatial-coherence observables — they probe phase variations across an aperture or between two separated telescopes — while $g_{\text{obs}}^{(1)}(\tau)$ derived above describes temporal coherence along a single line of sight. Nevertheless, both are governed by the same foam suppression factor $\mathcal{F}_{\text{foam}}$ of Eq. (24), because the underlying phase variance σ_ϕ^2 is common to the temporal and spatial projections of the full coherence function; this unifies the published constraints [4, 7–10] within a single framework. Where the temporal $g^{(1)}$ fails — quasi-static foam ($C_\alpha \approx 1$, no variation in time) — the spatial probes succeed; where the spatial probes fail — transversely coherent foam ($r_\perp \gg D_{\text{tel}}$, no variation across the aperture) — the temporal $g^{(1)}$ succeeds. The failure modes are complements because each probe samples the dimension the other does not. In the subsections below we adopt the temporal notation while noting the correspondence to the spatial channel.

1. Fringe Visibility and the Lieu–Hillman Condition

The fringe visibility measured in a stellar interferometer at time delay τ is $V(\tau) = |g_{\text{obs}}^{(1)}(\tau)|$. From Eq. (22):

$$V(\tau) = V_{\text{src}}(\tau) \cdot \mathcal{F}_{\text{foam}}(\tau), \quad (25)$$

where $V_{\text{src}}(\tau) = |g_{\text{src}}^{(1)}(\tau)|$ is the intrinsic source visibility. Lieu & Hillman [7] argue that fringes are erased when the accumulated phase randomizes over $[0, 2\pi)$, i.e. $\sigma_\phi \gtrsim 1$. In our language this is the limit $\mathcal{F}_{\text{foam}} \rightarrow 0$: the foam suppression factor collapses and $V \rightarrow 0$. The Lieu–Hillman condition $\Delta\phi \geq 2\pi$, where $\Delta\phi \simeq 2\pi\sigma_\ell(D)/\lambda$ [4], is precisely $\sigma_\phi \geq 1$. Our result extends this binary criterion to a continuous suppression curve $V(\tau)$ whose shape encodes the temporal decorrelation $C_\alpha(\tau; D)$.

2. Strehl Ratio

The Strehl ratio S — the ratio of the observed PSF peak to the diffraction-limited peak — is the standard measure of wavefront quality in telescope optics. In the Ng–Perlman picture [4], spacetime foam imprints a spatially varying random phase $\Delta\phi(x, y)$ across the entrance aperture: each point (x, y) on the mirror receives a ray that has accumulated an independent random path-length shift from the Planck-scale cells it traversed, producing a random phase screen with rms value σ_ϕ . When these phase shifts are spatially uncorrelated across the aperture — the limit $r_\perp \ll D_{\text{tel}}$, where $r_\perp$ is the transverse foam coherence length and D_{tel} is the aperture diameter — the Maréchal approximation gives

$$S = e^{-\sigma_\phi^2}, \quad (26)$$

valid for $\sigma_\phi \lesssim 1$; for $\sigma_\phi \gg 1$ the PSF peak is effectively destroyed.

Within the present temporal framework, the same quantity σ_ϕ^2 appears as the large-lag limit of the foam suppression factor. From Eq. (24), $\mathcal{F}_{\text{foam}}(\tau) = e^{-\sigma_\phi^2[1-C_\alpha(\tau; D)]}$, so in the dynamical limit $C_\alpha(\infty; D) \rightarrow 0$:

$$S = \mathcal{F}_{\text{foam}}(\tau \rightarrow \infty) = e^{-\sigma_\phi^2}. \quad (27)$$

This equality is an exact result of the Gaussian characteristic function, independent of the Maréchal approximation, but holds only in the dynamical limit $C_\alpha(\infty; D) \rightarrow 0$; for foam that does not fully decorrelate on any accessible timescale the two sides differ. The equality holds when both conditions are satisfied: $r_\perp \ll D_{\text{tel}}$ (so aperture elements sample independent foam realizations) and $C_\alpha(\infty; D) \rightarrow 0$ (so the single ray fully decorrelates at large lag); in either limit alone the two expressions differ. Both involve the same rms phase variance $\sigma_\phi^2 = (2\pi\sigma_\ell/\lambda)^2$. For quasi-static foam ($C_\alpha \approx 1$), $\mathcal{F}_{\text{foam}}(\infty) \rightarrow 1$ consistently with $S \rightarrow 1$ in the $r_\perp \gg D_{\text{tel}}$ limit discussed below.

It is important to distinguish the role of the transverse coherence function $\Phi(b/r_\perp)$ from the aperture-plane phase noise just described. The quantity σ_ϕ^2 is a single-ray variance: it characterises the accumulated phase fluctuation along one line of sight and does not by itself involve correlations between spatially separated rays. The Maréchal formula Eq. (26), however, additionally requires that adjacent aperture elements sample *independent* foam realizations, which is precisely the $r_\perp \ll D_{\text{tel}}$ condition already stated. In the opposite limit, $r_\perp \gg D_{\text{tel}}$, the entire aperture sits within a single coherence patch: the wavefront has a common phase offset that varies in time but not across the aperture, there is no spatial phase variation, and $S \rightarrow 1$ — the foam is invisible to PSF-based measurements even though σ_ϕ^2 may be large. The transverse coherence function $\Phi(b/r_\perp)$ governs a physically distinct effect — correlations between rays arriving from different sky directions $\theta \neq \theta'$ — which manifests in the second-order coherence of extended sources (Section VI) but does not directly enter the single-aperture Strehl ratio.

Both the Strehl and the Lieu–Hillman fringe visibility therefore measure the same underlying quantity $\sigma_\phi^2 = (2\pi\sigma_\ell/\lambda)^2$. The difference is observational channel: Strehl and fringes probe the *spatial* phase variation across the aperture (requiring $r_\perp \ll D_{\text{tel}}$), while the temporal $g^{(1)}$ of Eq. (22) probes the *temporal* phase variation along a single ray as a function of lag, with no assumption on $r_\perp$ relative to the aperture.

IV. SECOND-ORDER COHERENCE

We derive the second-order coherence $g_{\text{obs}}^{(2)}(\tau)$ for a point source, continuing the framework of Section III. The result has the same structure as Eq. (14): $g_{\text{obs}}^{(2)} - 1$ is a convolution of $g_{\text{src}}^{(2)} - 1$ with the identical Gaussian jitter kernel. Because the intensity $|E_{\text{obs}}|^2$ is blind to the carrier phase $e^{i\omega_0\delta t}$, foam leaves $g^{(2)}$ nearly unchanged in the narrow-band limit while strongly suppressing $g^{(1)}$ — the asymmetry that produces the Siegert violation developed in Section V.

A. Jitter Convolution Formula for $g^{(2)}$

Fix a realization of the jitter process $\delta t(\cdot)$. Conditioned on it, $E_{\text{obs}}(t)$ is a deterministic time-warp of a complex Gaussian field, so the Siegert identity holds at each fixed pair $(t, t + \tau)$ as a purely algebraic Gaussian moment identity, with no stationarity required. (Ergodicity of E_{src} allows the conditional time-average to be replaced by the conditional ensemble average; the subsequent average over the stationary δt process then yields a function of τ only, by independence of E_{src} and δt .) With $g_{\text{obs}}^{(1)}(\tau | \delta t) = g_{\text{src}}^{(1)}(\tau - \Delta(\tau))$ and the intrinsic Siegert relation $g_{\text{src}}^{(2)} - 1 = |g_{\text{src}}^{(1)}|^2$:

$$g_{\text{obs}}^{(2)}(\tau | \delta t) - 1 = g_{\text{src}}^{(2)}(\tau - \Delta(\tau)) - 1. \quad (28)$$

Averaging over the jitter distribution:

$$g_{\text{obs}}^{(2)}(\tau) - 1 = \langle g_{\text{src}}^{(2)}(\tau - \Delta(\tau)) - 1 \rangle_{\delta t} = \int_{-\infty}^{\infty} \mathcal{N}(\Delta; 0, \sigma_{\Delta}^2(\tau)) [g_{\text{src}}^{(2)}(\tau - \Delta) - 1] d\Delta. \quad (29)$$

This Gaussian kernel relies on the CLT argument of Section II C, which requires $\tau \gg \tau_{\text{foam}}$; for $\tau \lesssim \tau_{\text{foam}}$ the kernel may deviate from Gaussian and Eq. (29) should be regarded as approximate in that regime (see Section II B for the definition and limits of τ_{foam}).

B. Narrow-Band Limit: $g^{(2)}$ is Unchanged to Leading Order

The intensity $I_{\text{obs}}(t) = |E_{\text{obs}}(t)|^2$ is blind to the carrier phase $e^{i\omega_0 \delta t(t)}$: only the envelope matters. In the narrow-band limit $\delta t \ll \tau_c$, the jitter does not appreciably shift the envelope, $A(t - \delta t) \approx A(t)$, so $I_{\text{obs}}(t) \approx |A(t)|^2 = I_{\text{src}}(t)$ to leading order in $\delta t/\tau_c$. Consequently:

$$g_{\text{obs}}^{(2)}(\tau) = g_{\text{src}}^{(2)}(\tau) \quad (\text{leading order in } \delta t/\tau_c), \quad (30)$$

with corrections of order $(\sigma_{\ell}/c\tau_c)^2$ from both envelope-timing and carrier-phase contributions to the intensity. For $\alpha = 1/2$ at $D = 1$ Gpc, $\lambda = 500$ nm, $\Delta\lambda/\lambda = 10^{-3}$: $\delta t_{\text{rms}}/\tau_c \approx 0.28$, so corrections are at the $\sim 8\%$ level and the exact convolution Eq. (29) should be used for precision work. For $\alpha = 2/3$ the expansion parameter is $\sim 10^{-11}$ and Eq. (30) is exact to any practical precision.

By contrast, $g^{(1)}$ is sensitive to the carrier phase because $\sigma_{\phi} = \omega_0 \sigma_{\ell}/c$ can be large even when $\delta t \ll \tau_c$: for the same parameters, $\sigma_{\phi} \approx 280 \gg 1$.

V. SIEGERT RELATION VIOLATION

The Siegert relation $g^{(2)}(\tau) - 1 = |g^{(1)}(\tau)|^2$ holds for any stationary complex Gaussian field and provides a powerful internal consistency check: once $g^{(1)}$ is known, the expected $g^{(2)}$ is fixed. As foreshadowed in Section IV, foam violates this relation because $g_{\text{obs}}^{(1)}$ is suppressed by the carrier-phase factor $\mathcal{F}_{\text{foam}}$ while $g_{\text{obs}}^{(2)}$ is not. This section derives and characterizes the Siegert discrepancy $\mathcal{S}(\tau) \equiv [g_{\text{obs}}^{(2)} - 1] - |g_{\text{obs}}^{(1)}|^2$.

A. General Bound

Define the Siegert discrepancy:

$$\mathcal{S}(\tau) \equiv [g_{\text{obs}}^{(2)}(\tau) - 1] - |g_{\text{obs}}^{(1)}(\tau)|^2. \quad (31)$$

Let $X \equiv g_{\text{src}}^{(1)}(\tau - \Delta(\tau))$, a complex random variable. From Eqs. (14) and (29):

$$g_{\text{obs}}^{(1)}(\tau) = \langle X \rangle_{\delta t}, \quad g_{\text{obs}}^{(2)}(\tau) - 1 = \langle |X|^2 \rangle_{\delta t}. \quad (32)$$

Therefore

$$\mathcal{S}(\tau) = \langle |X|^2 \rangle - |\langle X \rangle|^2 \geq 0, \quad (33)$$

by the Cauchy–Schwarz inequality, with equality only when $\Delta(\tau)$ has zero variance (i.e., at $\tau = 0$ or in the absence of foam). The Siegert relation is violated for all $\tau \neq 0$ whenever foam is present. This non-negativity holds for a point source; for extended sources the cross-ray differential jitter can drive $\mathcal{S}(0) < 0$ (see Section VI). The Siegert violation is not a non-classical signature: the foam-modified field remains a classical stochastic process, and the classical inequalities $g^{(2)}(0) \geq g^{(2)}(\tau)$ and $g^{(2)}(0) > 1$ are preserved; it is only the Gaussian-statistics identity relating $g^{(2)}$ to $g^{(1)}$ that is broken, due to the stochastic phase accumulated along the propagation path.

B. Sensitivity Scale Decoupling

The violation arises because $g^{(1)}$ and $g^{(2)}$ respond to foam on different scales. The first-order coherence is controlled by $\sigma_\phi = 2\pi\sigma_\ell(D)/\lambda$, the rms accumulated phase along a single ray; $g^{(1)}$ is sensitive to it because the carrier phase $e^{i\omega_0\delta t}$ appears directly in the field correlation. The second-order coherence is controlled by $\sigma_\Delta(\tau)/\tau_c$, the rms differential jitter scaled by the coherence time; $g^{(2)}$ is sensitive to it because intensity correlations depend on how much the jitter *changes* between two detection events, not its absolute value. These two scales are related by

$$\frac{\sigma_\Delta(\tau)}{\tau_c} = \frac{\sqrt{2[1 - C_\alpha(\tau)]} \sigma_\ell(D)}{c\tau_c} = \sigma_\phi \cdot \frac{\Delta\lambda}{\lambda} \cdot \sqrt{2[1 - C_\alpha(\tau)]}, \quad (34)$$

using $\omega_0\tau_c = \lambda/\Delta\lambda$. The factor $\Delta\lambda/\lambda \ll 1$ means that even when $\sigma_\phi \gg 1$ (fringes erased), $\sigma_\Delta(\tau)/\tau_c$ can be small — so $g^{(2)}$ remains near its intrinsic value while $g^{(1)}$ is suppressed. This decoupling is the physical root of the Siegert violation. Numerical values of both scales are given in Table I (evaluated at the dynamical limit $C_\alpha \rightarrow 0$; for quasi-static foam $C_\alpha \approx 1$ the ratio vanishes at all lags).

TABLE I. Sensitivity scales at $D = 1$ Gpc, $\lambda = 500$ nm, $\Delta\lambda/\lambda = 10^{-3}$, $A_\alpha = 1$, $C_\alpha \rightarrow 0$. The $\alpha = 1/2$ case is excluded by imaging under the standard assumption of short-range transverse foam correlations [10]; models with long-range correlations evade this bound [11]. $\alpha = 2/3$ (holographic) is the most sensitive surviving model.

Scale	$\alpha = 1/2$	$\alpha = 2/3$
σ_ϕ	280	2.5×10^{-8}
$\sigma_\Delta(\tau)/\tau_c$	0.40	3.6×10^{-11}

C. Explicit Form and Properties in the Narrow-Band Limit

In the narrow-band limit, $g_{\text{obs}}^{(2)} \approx g_{\text{src}}^{(2)}$ (Eq. (30)), while $g_{\text{obs}}^{(1)}$ is suppressed by $\mathcal{F}_{\text{foam}}(\tau)$ (Eq. (22)). Using the intrinsic Siegert relation $g_{\text{src}}^{(2)}(\tau) - 1 = |g_{\text{src}}^{(1)}(\tau)|^2$, to leading order:

$$\mathcal{S}(\tau) \approx |g_{\text{src}}^{(1)}(\tau)|^2 \left(1 - e^{-2\sigma_\phi^2[1 - C_\alpha(\tau; D)]} \right). \quad (35)$$

In the weak-jitter limit $\sigma_\phi \ll 1$, $\mathcal{S} \approx 2\sigma_\phi^2[1 - C_\alpha] |g_{\text{src}}^{(1)}|^2$ — linear in σ_ϕ^2 and small. In the strong-jitter limit $\sigma_\phi \gg 1$, the exponential vanishes and $\mathcal{S} \approx |g_{\text{src}}^{(1)}(\tau)|^2$: $g_{\text{obs}}^{(2)} - 1$ retains its intrinsic value while $|g_{\text{obs}}^{(1)}|^2 \rightarrow 0$, producing a full violation. Quasi-static foam ($C_\alpha \approx 1$) gives $\mathcal{S} \approx 0$ at all lags — foam is invisible to all coherence observables in this limit. The discrepancy $\mathcal{S}(\tau)$ is a direct observable — both $g_{\text{obs}}^{(1)}$ and $g_{\text{obs}}^{(2)}$ are measurable from the same dataset — and the modulation beyond the prefactor $|g_{\text{src}}^{(1)}(\tau)|^2$ is controlled solely by α , D , and $C_\alpha(\tau; D)$, independently of the unknown source spectrum. For a Lorentzian source, $|g_{\text{src}}^{(1)}(\tau_c)|^2 = e^{-2}$, and in the dynamical limit $C_\alpha \rightarrow 0$, Eq. (35) reduces to

$$\mathcal{S}(\tau_c) = e^{-2} \left(1 - e^{-2\sigma_\phi^2} \right), \quad (36)$$

the detection observable at lag τ_c . In the weak-foam limit $\sigma_\phi \ll 1$ this is $\approx 2e^{-2}\sigma_\phi^2$; in the strong-foam limit $\sigma_\phi \gg 1$ it saturates at $e^{-2} \approx 0.135$. Although derived here for a point source, Eq. (35) applies equally to extended sources at finite lag $\tau \neq 0$: $g_{\text{src}}^{(1)}(\tau)$ is unchanged (Section VI), and the additional cross-ray jitter that modifies $g_{\text{obs}}^{(2)}$ at $\tau = 0$ produces only $O((\sigma_{\text{eff}}/\tau_c)^2)$ corrections at finite lag in the narrow-band limit. The extended-source novelty — the linear-in- s bunching suppression $\delta g^{(2)}(0)$ — is confined to zero lag, where the cusp in $g_{\text{src}}^{(2)}$ amplifies the cross-ray differential jitter into a first-order effect.

VI. EXTENDED SOURCES: TRANSVERSE FOAM COHERENCE

We now generalize from a point source to an extended incoherent source. The first-order coherence $g_{\text{obs}}^{(1)}$ turns out to be unchanged from the point-source result, because $g^{(1)}$ probes only single-ray statistics. The second-order coherence $g_{\text{obs}}^{(2)}$, however, acquires new structure through cross-ray differential jitter between rays arriving from different sky directions, producing a suppression of the zero-lag bunching peak that directly probes the transverse foam coherence length $r_{\perp}$.

A. Field Model and Spatiotemporal Covariance

We generalize to an extended incoherent source whose angular brightness distribution is $I(\boldsymbol{\theta})$ (normalized to unity). The analysis assumes a point (single spatial mode) detector; a finite aperture would introduce Van Cittert–Zernike geometric phases that modify the angular coherence. The detected field is a weighted superposition over rays labelled by direction offset $\boldsymbol{\theta}$:

$$E_{\text{obs}}(t) = \int d\boldsymbol{\theta} I^{1/2}(\boldsymbol{\theta}) E_{\boldsymbol{\theta}}(t - \delta t(t; \boldsymbol{\theta})), \quad (37)$$

where the $E_{\boldsymbol{\theta}}$ are independent, identically distributed complex Gaussian processes (one per sky direction), and each ray carries its own foam jitter $\delta t(t; \boldsymbol{\theta})$. The spatiotemporal covariance of the path-length aggregate is given by Eq. (5) of Section II, with $b(\boldsymbol{\theta}, \boldsymbol{\theta}') \simeq D|\boldsymbol{\theta} - \boldsymbol{\theta}'|$ the transverse separation between rays evaluated at the source end. For foam distributed uniformly along the line of sight the ray separation at position z is $z|\boldsymbol{\theta} - \boldsymbol{\theta}'|$, so the full cross-covariance integral involves an average over $0 \leq z \leq D$; setting $b = D|\boldsymbol{\theta} - \boldsymbol{\theta}'|$ evaluates the separation at the source plane. Because $r_{\perp}$ is treated as an effective free parameter, the $O(1)$ geometric factor from the z -integral is absorbed into its definition; the angular-size dependence of $\delta g^{(2)}$ derived in Section VID is the corresponding observational signature.

B. First-Order Coherence is Unchanged

The numerator of $g_{\text{obs}}^{(1)}(\tau)$ involves cross terms $\langle E_{\boldsymbol{\theta}}^*(\dots) E_{\boldsymbol{\theta}'}(\dots) \rangle$. Because different source directions are spatially incoherent,

$$\langle E_{\boldsymbol{\theta}}^*(s) E_{\boldsymbol{\theta}'}(s') \rangle \propto \delta(\boldsymbol{\theta} - \boldsymbol{\theta}'), \quad (38)$$

so only the diagonal $\boldsymbol{\theta} = \boldsymbol{\theta}'$ contributes. Each diagonal term reproduces the single-ray result of Section III C, and the angular integral over I gives unity. Therefore:

$$g_{\text{obs}}^{(1)}(\tau) = g_{\text{src}}^{(1)}(\tau) \cdot e^{-\sigma_{\phi}^2[1 - C_{\alpha}(\tau; D)]}, \quad (39)$$

identical to the point-source result Eq. (22). The first-order coherence probes only single-ray statistics; the transverse foam structure leaves no imprint on $g^{(1)}$.

C. Second-Order Coherence for an Extended Source

The unnormalized intensity correlator is $\langle I(t)I(t+\tau) \rangle = \langle E^*(t)E(t)E^*(t+\tau)E(t+\tau) \rangle$, where $E \equiv E_{\text{obs}}$ is the field of Eq. (37). Expanding the double angular integral and applying the Wick (moment) theorem to the four-point function of the zero-mean complex Gaussian field $E_{\boldsymbol{\theta}}$ produces two pairings (the $I^{1/2}$ weight factors enter the angular integral and are suppressed here for clarity):

$$\langle E_{\boldsymbol{\theta}_1}^* E_{\boldsymbol{\theta}_2} E_{\boldsymbol{\theta}_3}^* E_{\boldsymbol{\theta}_4} \rangle = \langle E_{\boldsymbol{\theta}_1}^* E_{\boldsymbol{\theta}_2} \rangle \langle E_{\boldsymbol{\theta}_3}^* E_{\boldsymbol{\theta}_4} \rangle + \langle E_{\boldsymbol{\theta}_1}^* E_{\boldsymbol{\theta}_4} \rangle \langle E_{\boldsymbol{\theta}_3}^* E_{\boldsymbol{\theta}_2} \rangle. \quad (40)$$

The first pairing contributes $\langle I(t) \rangle \langle I(t+\tau) \rangle$ (the disconnected term, which cancels when forming $g^{(2)} - 1$). The second pairing gives the connected contribution. Because different source directions are spatially incoherent, $\langle E_{\boldsymbol{\theta}}^*(s) E_{\boldsymbol{\theta}'}(s') \rangle \propto \delta(\boldsymbol{\theta} - \boldsymbol{\theta}')$, so the connected pairing requires $\boldsymbol{\theta}_1 = \boldsymbol{\theta}_4$ and $\boldsymbol{\theta}_3 = \boldsymbol{\theta}_2$. The four-fold angular integral therefore collapses to a double integral over ray pairs $(\boldsymbol{\theta}, \boldsymbol{\theta}')$, each pair contributing its own jitter-convolved $g_{\text{src}}^{(2)}$:

$$g_{\text{obs}}^{(2)}(\tau) - 1 = \iint d\boldsymbol{\theta} d\boldsymbol{\theta}' I(\boldsymbol{\theta}) I(\boldsymbol{\theta}') \int_{-\infty}^{\infty} \mathcal{N}(\Delta; 0, \sigma_{\Delta}^2(\tau; \boldsymbol{\theta}, \boldsymbol{\theta}')) [g_{\text{src}}^{(2)}(\tau - \Delta) - 1] d\Delta, \quad (41)$$

where $\sigma_\Delta^2(\tau; \boldsymbol{\theta}, \boldsymbol{\theta}')$ is from Eq. (8) (Section II C). The key new feature relative to the point-source case is that for $b \gg r_\perp$, $\Phi \rightarrow 0$ and $\sigma_\Delta^2 \rightarrow 2\sigma_\ell^2/c^2$ at all lags, including $\tau = 0$. In the limit $\Phi \equiv 1$ (all rays transversely coherent), every kernel reduces to the single-ray $\mathcal{N}(0, \sigma_\Delta^2(\tau))$ and Eq. (41) reduces to the point-source result Eq. (29).

D. Effective Transverse Coherence Factor

To reduce Eq. (41) to a tractable form, define the effective transverse coherence factor:

$$\Phi_{\text{eff}} \equiv \iint d\boldsymbol{\theta} d\boldsymbol{\theta}' I(\boldsymbol{\theta}) I(\boldsymbol{\theta}') \Phi\left(\frac{b(\boldsymbol{\theta}, \boldsymbol{\theta}')}{r_\perp}\right). \quad (42)$$

This replaces the ray-dependent differential variance with a single effective variance:

$$\sigma_{\tau, \text{eff}}^2(\tau) \equiv \frac{2\sigma_\ell^2(D)}{c^2} \left[1 - C_\alpha(\tau; D) \Phi_{\text{eff}}\right], \quad (43)$$

and allows Eq. (41) to be approximated by the simpler 1D convolution of Eq. (29) with σ_Δ^2 replaced by $\sigma_{\tau, \text{eff}}^2$. The two limits are: $\Phi_{\text{eff}} \rightarrow 1$ (point-like source or $r_\perp \gg D \Delta\theta$) recovers the single-ray result; $\Phi_{\text{eff}} \rightarrow 0$ (fully incoherent, $r_\perp \ll D \Delta\theta$) maximizes the effective variance. The approximation becomes exact in both limits. At intermediate angular sizes ($r_\perp \sim D \Delta\theta$), $\Phi(b/r_\perp)$ varies across ray pairs and the Φ_{eff} reduction introduces errors that have not been quantified here; they must be assessed jointly with the CLT breakdown of Section II C, since both occur in the same regime $b \sim r_\perp$, and we leave a worked numerical comparison to future work.

E. Extended-Source Siegert Discrepancy and Zero-Lag Bunching

For the extended source, $g_{\text{obs}}^{(1)}(\tau)$ is unchanged from the point-source result (Eq. 39), while $g_{\text{obs}}^{(2)}(\tau)$ is given by the general mixture Eq. (41). Unlike the point-source case, in the narrow-band limit $g_{\text{obs}}^{(2)}$ is *not* equal to $g_{\text{src}}^{(2)}$: the cross-ray differential jitter modifies the intensity correlation through the envelope timing, even though the carrier phase does not affect individual intensities. The Siegert discrepancy is therefore

$$\mathcal{S}(\tau) = \left[g_{\text{obs}}^{(2)}(\tau) - 1\right] - \left|g_{\text{obs}}^{(1)}(\tau)\right|^2, \quad (44)$$

where $g_{\text{obs}}^{(2)}(\tau) - 1$ is given by Eq. (41) and $|g_{\text{obs}}^{(1)}(\tau)|^2 = |g_{\text{src}}^{(1)}(\tau)|^2 e^{-2\sigma_\phi^2[1 - C_\alpha(\tau)]}$. At finite lag the cross-ray jitter correction to $g_{\text{obs}}^{(2)}$ is $O((\sigma_{\text{eff}}/\tau_c)^2)$ in the narrow-band limit, and $\mathcal{S}(\tau \neq 0)$ reduces to the point-source result Eq. (35); the remainder of this subsection focuses on $\tau = 0$, where the cusp produces a qualitatively distinct signature.

At $\tau = 0$, $g_{\text{obs}}^{(1)}(0) = 1$ exactly by the spatial-incoherence argument of Section VI, so $\mathcal{S}(0) = g_{\text{obs}}^{(2)}(0) - 2$. Evaluating the zero-lag effective variance with $C_\alpha(0; D) = 1$ gives

$$\sigma_0^2 \equiv \frac{2\sigma_\ell^2(D)}{c^2} (1 - \Phi_{\text{eff}}), \quad (45)$$

which is strictly positive whenever $\Phi_{\text{eff}} < 1$, so the ray-pair mixture reduces to a single Gaussian kernel of width σ_0 .

For a Lorentzian source, $g_{\text{src}}^{(1)}(\tau) = e^{-|\tau|/\tau_c}$, so the Siegert relation gives $g_{\text{src}}^{(2)}(\tau) - 1 = e^{-2|\tau|/\tau_c}$. The zero-lag convolution over $\Delta \sim \mathcal{N}(0, \sigma_0^2)$ then becomes, setting $s \equiv \sigma_0/\tau_c$ and splitting on the sign of Δ (both halves contributing equally by symmetry of $|\Delta|$),

$$\langle e^{-2|\Delta|/\tau_c} \rangle = \frac{2}{\sqrt{2\pi} \sigma_0} \int_0^\infty e^{-2\Delta/\tau_c} e^{-\Delta^2/2\sigma_0^2} d\Delta = e^{2s^2} \text{erfc}(\sqrt{2}s), \quad (46)$$

which equals 1 as $s \rightarrow 0$ (no suppression) and 0 as $s \rightarrow \infty$ via the standard asymptotic $e^{x^2} \text{erfc}(x) \sim 1/(x\sqrt{\pi})$ (full suppression, $g_{\text{obs}}^{(2)}(0) \rightarrow 1$). So

$$g_{\text{obs}}^{(2)}(0) = 1 + e^{2s^2} \text{erfc}(\sqrt{2}s), \quad s \equiv \sigma_0/\tau_c, \quad (47)$$

and the zero-lag bunching peak is suppressed below its ideal value of 2 by

$$\delta g^{(2)} \equiv 2 - g_{\text{obs}}^{(2)}(0) = 1 - e^{2s^2} \text{erfc}(\sqrt{2}s), \quad (48)$$

In the small- s limit, $\delta g^{(2)} \approx 2\sqrt{2/\pi} s + O(s^2)$, so the suppression is linear in s for a Lorentzian source.

$$\mathcal{S}(0) = -\delta g^{(2)} \leq 0. \quad (49)$$

This sign is qualitatively distinct from the point-source result $\mathcal{S}(\tau \neq 0) \geq 0$ of Section V: a negative zero-lag Siegert discrepancy is therefore a sign-flip diagnostic that distinguishes spatially resolved from unresolved foam signatures. In the point-source limit $\Phi_{\text{eff}} \rightarrow 1$, $\sigma_0 \rightarrow 0$, $\delta g^{(2)} \rightarrow 0$, and $\mathcal{S}(0) \rightarrow 0$, recovering the point-source result. Since $\sigma_0^2 \propto (1 - \Phi_{\text{eff}})$ and Φ_{eff} decreases as the source angular diameter increases at fixed D and $r_\perp$, more extended sources experience greater suppression of the bunching peak and a more negative Siegert discrepancy at zero lag.

For an astrophysically more representative Voigt profile — a Lorentzian of width τ_L convolved with a Gaussian of width τ_G (modelling natural plus Doppler/turbulent broadening) — the coherence function is $g_{\text{src}}^{(1)}(\tau) = e^{-|\tau|/\tau_L} e^{-\tau^2/2\tau_G^2}$, giving $g_{\text{src}}^{(2)}(\tau) - 1 = e^{-2|\tau|/\tau_L} e^{-\tau^2/\tau_G^2}$. The zero-lag convolution over $\Delta \sim \mathcal{N}(0, \sigma_0^2)$ combines the source Gaussian with the jitter kernel into an effective width

$$\tilde{\sigma}^2 = \frac{\sigma_0^2 \tau_G^2}{2\sigma_0^2 + \tau_G^2}, \quad (50)$$

yielding

$$g_{\text{obs}}^{(2)}(0) - 1 = \frac{\tilde{\sigma}}{\sigma_0} e^{2\tilde{s}^2} \text{erfc}(\sqrt{2}\tilde{s}), \quad \tilde{s} \equiv \tilde{\sigma}/\tau_L. \quad (51)$$

For any Voigt profile with finite τ_L , the parameter $\tilde{s} > 0$ and an $O(s)$ term is present, controlled by the Lorentzian component. A purely Gaussian source ($\tau_L \rightarrow \infty$, $\tilde{s} \rightarrow 0$) has no $O(s)$ correction: $g_{\text{obs}}^{(2)}(0) - 1 = 1/\sqrt{1 + 2s^2(\tau_c/\tau_G)^2}$.

In the narrow-band limit $\sigma_\ell/(c\tau_c) \ll 1$ (Section III C), $\sigma_0 \ll \tau_c$ for any Φ_{eff} . Substituting $u = \Delta/\sigma_0$ and expanding $g_{\text{src}}^{(2)}(\sigma_0 u) - 1$ for small σ_0 , the convolution integral yields to first order in $s = \sigma_0/\tau_c$:

$$g_{\text{obs}}^{(2)}(0) - 1 \approx [g_{\text{src}}^{(2)}(0) - 1] + \sqrt{\frac{2}{\pi}} \sigma_0 g_{\text{src}}^{(2)'}(0^+), \quad (52)$$

where $g_{\text{src}}^{(2)'}(0^+)$ denotes the right-derivative of $g_{\text{src}}^{(2)}$ at zero lag (the function is non-differentiable at $\tau = 0$ for cusped profiles such as the Lorentzian, so the symmetric derivative does not exist there). For sources with a smooth (even) profile the right-derivative vanishes and the leading correction is $O(\sigma_0^2)$; for a Lorentzian, $g_{\text{src}}^{(2)'}(0^+) = -2/\tau_c$, recovering $\delta g^{(2)} \approx 2\sqrt{2/\pi} s$, consistent with the small- s expansion of Eq. (48). For a Voigt profile, the same expansion gives $\delta g^{(2)} \approx 2\sqrt{2/\pi} \sigma_0/\tau_L + (\sigma_0/\tau_G)^2 + \dots$: the $O(s)$ term is controlled solely by the Lorentzian component τ_L ; maximizing foam sensitivity therefore favours spectrally resolved emission lines with a Lorentzian component, since an unresolved line or a pure continuum erases the cusp and reverts to $O(s^2)$.

The suppression $\delta g^{(2)}$ arises from a purely spatial mechanism with no analogue in the point-source case: even at $\tau = 0$, photon pairs from different sky directions $\boldsymbol{\theta} \neq \boldsymbol{\theta}'$ accumulate independent instantaneous jitters $\delta t(t; \boldsymbol{\theta}) \neq \delta t(t; \boldsymbol{\theta}')$ when $b \gg r_\perp$, and their relative offset blurs the $g^{(2)}(0)$ measurement. Because $|g_{\text{obs}}^{(1)}(0)|^2 = 1$ exactly is unaffected by this spatial jitter, the two observables are controlled by different mechanisms and $\mathcal{S}(0) < 0$. This does not violate any fundamental inequality: the bound $\mathcal{S} \geq 0$ established in Section V applies to a single ray, and the angular mixture of an extended source need not satisfy it.

For the region $\alpha > 1/2$ that experiments seek to probe, $\sigma_\phi \ll 1$ and the sensitivity parameter $s = \sigma_0/\tau_c = \sqrt{2(1 - \Phi_{\text{eff}})} \sigma_\phi \cdot (\Delta\lambda/\lambda)$ is small across all realistic observational bandwidths. For the holographic model ($\alpha = 2/3$), $\sigma_\phi \sim 10^{-8}$, so that even at a high spectral resolution $R = \lambda/\Delta\lambda \sim 10^4$, one finds $s \sim 10^{-12}$. In this regime the linear approximation $\delta g^{(2)} \approx 2\sqrt{2/\pi} s$ is essentially exact.

Beyond raw detectability, the coherence observables serve as discriminants against conventional propagation phase screens [3]. Furthermore, for foam models with long-range transverse correlations image blurring is strongly suppressed [11], making the coherence observables developed here potentially the only channel sensitive to such models. The source-size dependence of $\delta g^{(2)}$ through Φ_{eff} provides a direct probe of $r_\perp$ with no analogue in a featureless propagation screen, which carries no transverse coherence structure. The mathematical structure of the foam suppression factor $\mathcal{F}_{\text{foam}}$ is identical to that of any Gaussian random phase screen, but two signatures distinguish foam: distance scaling and achromaticity. For spacetime foam $\sigma_\phi^2 \propto D^{2(1-\alpha)}$ (D^1 for $\alpha = 1/2$, $D^{2/3}$ for $\alpha = 2/3$, D^0 for $\alpha = 1$), whereas a thin intervening screen scales as D and Kolmogorov turbulence as $D^{11/6}$; the holographic and conservative foam models predict sub-linear or flat growth qualitatively different from any known astrophysical screen. Geometrically, foam jitter ($\delta t = \Delta\ell/c$) is achromatic — independent of λ — in contrast to plasma-based effects which scale as λ^2 . The Fermi, H.E.S.S., and LHAASO time-of-flight analyses [12–14] rule out deterministic, energy-ordered LIV

dispersion as a class of quantum-gravity effect, but do not constrain the stochastic achromatic foam jitter considered here (Section I); the principal chromatic alternative that the achromaticity discriminant addresses is therefore plasma propagation. In practice, the intrinsic source visibility $V_{\text{src}}(\tau)$ introduces a degeneracy with foam suppression, but the Siegert discrepancy $\mathcal{S}(\tau) = g_{\text{obs}}^{(2)}(\tau) - 1 - |g_{\text{obs}}^{(1)}(\tau)|^2$ is constructed directly from the data: a detection of $\mathcal{S} \neq 0$ requires no explicit model of the source spectrum, since $g_{\text{src}}^{(1)}(\tau)$ can be obtained from the source spectrum via the Wiener–Khinchin theorem.

VII. SIEGERT-VIOLATION: SENSITIVITY AND TARGET SELECTION

This section examines the detectability of Siegert violation via the HBT effect and surveys candidate astrophysical sources: broad-line AGN and quasars, Galactic and Local-Group transients, and nearby stars. We use the Hanbury Brown–Twiss (HBT) split-beam configuration [16] throughout, as essentially all modern photon-counting $g^{(2)}$ measurements do: detector dead time makes a single-detector autocorrelation blind to the ps-scale lags that build the bunching peak. The sensitivity estimates that follow assume a Lorentzian line profile, adopted as the default for the broad and narrow permitted lines considered here. Two variants are considered: the temporal decorrelation/non-zero lag; transverse decorrelation/zero lag. The pure-Lorentzian cusp in $g_{\text{src}}^{(2)}$ at $\tau = 0$ maximizes the linear-in- s scaling of Eq. (48); any departure from a pure Lorentzian rounds the cusp and reverts the extended-source discrepancy to $O(s^2)$, so the SNR estimates for that variant should be read as optimistic limits. Filter strategy differs between the finite-lag ($\tau \neq 0$) and zero-lag ($\tau = 0$) variants and is discussed in each subsubsection in turn. All these targets are physically large and are thus extended sources, even though they are difficult to resolve from Earth; as extended sources, they are sensitive to spacetime foam whose transverse coherence length satisfies $r_{\perp} \lesssim D \Delta\theta$ (equivalently $\Phi_{\text{eff}} \approx 0$), through the cross-ray differential jitter mechanism.

A. Sensitivity and Aperture Scaling

The HBT observable is $g_{\text{obs}}^{(2)}(\tau)$, whose foam-free value for a thermal source equals 2 at $\tau = 0$ and decays to 1 on the source coherence timescale. Foam imprints two distinct signatures on this observable: a Siegert discrepancy at finite lag (Section V), applicable to both point and extended sources, and a suppression of the zero-lag bunching peak specific to spatially extended sources (Section VI). Both are inferred by comparing the measured $g_{\text{obs}}^{(2)}$ to the Siegert prediction $g_{\text{src}}^{(2)} - 1 = |g_{\text{src}}^{(1)}|^2$, built from the source spectrum via the Wiener–Khinchin theorem for the case of $\tau \neq 0$; neither requires a separate model of the source. A detection at n - σ requires a signal uncertainty $\sigma[g_{\text{obs}}^{(2)}] \leq |\text{discrepancy}|/n$, where the discrepancy amplitude and the relevant lag depend on which variant is being measured.

The photon-counting rate is $R = \epsilon A \dot{n}$, where $\dot{n}$ is the source photon flux per unit area within the filter bandpass, ϵ the end-to-end throughput, and A the collecting area. We adopt the baseline instrumental configuration of Kim *et al.* [20]: telescope diameter $d = 9.96$ m ($A = 77.9$ m²), mirror reflectivity 0.73, filter transmission 0.90, detector quantum efficiency 0.60, total throughput $\epsilon = 0.39$, and timing resolution FWHM = 30 ps ($\sigma_t = 13$ ps rms). The astrophysical targets of interest have intrinsic coherence times $\tau_L \lesssim 1$ ps, so $\sigma_t \gg \tau_L$ for the baseline instrument and the measurement sits firmly in the detector-limited regime: an unresolved foam feature of width τ_L (or τ_c) is diluted to $(\tau_L/\sigma_t) |\mathcal{S}|$ in a time bin of width σ_t , with statistical noise $\sigma[g_{\text{meas}}^{(2)}] = 1/(R\sqrt{\sigma_t T_{\text{obs}}})$. We present the detector-limited SNR as the primary case throughout and quote the source-limited limit ($\tau_L \gtrsim \sigma_t$, requiring sub-ps timing) only as the gain that next-generation detectors would unlock with precision optics. The two variants differ in their observable, their signal scaling with the foam amplitude, and their distance and aperture dependence; we treat each in turn.

1. Temporal Decorrelation – Finite-lag variant

For a point-like source, the foam signature manifests as a Siegert discrepancy at finite lag $\tau \neq 0$. Considering a Lorentzian source at the dynamical limit $C_{\alpha}(\tau_c; D) \rightarrow 0$, the discrepancy at $\tau = \tau \sim \tau_c$ is given by Eq. 36:

$$\mathcal{S}(\tau_c) = e^{-2} \left(1 - e^{-2\sigma_{\phi}^2} \right), \quad (53)$$

where $\sigma_{\phi} = 2\pi\sigma_{\ell}(D)/\lambda$ is the rms accumulated phase (Eq. 23). The signal is *quadratic* in σ_{ϕ} in the weak-foam limit, $\mathcal{S}(\tau_c) \approx 2e^{-2}\sigma_{\phi}^2$, and saturates at $e^{-2} \approx 0.135$ for $\sigma_{\phi} \gg 1$.

a. Detector-limited SNR. With $\sigma_t \gg \tau_c$ for all sources of interest, the bunching feature of width τ_c is unresolved in a time bin of width σ_t , and

$$\text{SNR}_{\text{det}}^{\text{pt}} = \frac{\tau_c}{\sigma_t} |\mathcal{S}(\tau_c)| \cdot R \sqrt{\sigma_t T_{\text{obs}}} \xrightarrow{\sigma_\phi \ll 1} 2e^{-2} \sigma_\phi^2 R \frac{\tau_c}{\sqrt{\sigma_t}} \sqrt{T_{\text{obs}}}. \quad (54)$$

The detector-limited form carries a $\sqrt{\tau_c/\sigma_t}$ penalty relative to source-limited operation ($\text{SNR}_{\text{src}}^{\text{pt}} = |\mathcal{S}(\tau_c)| R \sqrt{\tau_c T_{\text{obs}}}$); reaching the source-limited limit on broad-line quasars requires $\sigma_t \lesssim \tau_c \sim 0.1$ ps.

b. Filter strategy. Unlike the zero-lag variant, preserving the intrinsic line cusp is not required: the finite-lag Siegert discrepancy $\mathcal{S}(\tau_c)$ depends only on σ_ϕ , not on the spectral shape that sets the coherence time. Filtering narrower than the intrinsic linewidth is therefore permitted, with τ_c replaced by the post-filter coherence time $\tau_{c,\text{filt}} = \lambda^2/(c \Delta \lambda_{\text{filter}})$. In the detector-limited regime the SNR scales as $R \tau_c / \sqrt{\sigma_t}$ (Eq. 54); for a Lorentzian emission line the transmitted fraction scales as $\Delta \lambda_{\text{filter}} / \Delta \lambda_{\text{line}}$ in the narrow-filter limit, so $R \tau_c$ is approximately constant and the SNR is insensitive to filter width. The same holds for a flat continuum source where $R \propto \Delta \lambda_{\text{filter}}$ and $\tau_c \propto \Delta \lambda_{\text{filter}}^{-1}$. The practical bound on filter width is set by sky background: widening the filter beyond the emission line adds sky photons without increasing source signal, degrading the signal-to-background ratio. For continuum sources the sky background constrains the maximum useful $\Delta \lambda_{\text{filter}}$; the optimal width is therefore set by the sky-limited signal-to-background ratio rather than by coherence considerations.

c. Distance and aperture scaling. Substituting $R \propto D^{-2}$ at fixed luminosity and $\sigma_\phi^2 \propto D^{2(1-\alpha)}$ into Eq. 54 gives, in the weak-foam limit,

$$\text{SNR}^{\text{pt}} \propto D^{2(1-\alpha)} \cdot D^{-2} = D^{-2\alpha}, \quad (55)$$

saturating at $\text{SNR} \propto D^{-2}$ in the strong-foam limit. Since $\alpha > 0$, the SNR increases toward smaller D : nearby targets are preferred.

The minimum detectable path-length amplitude follows from inverting Eq. 54 with the quadratic signal dependence $|\mathcal{S}(\tau_c)| \propto \sigma_\phi^2 \propto \sigma_\ell^2$:

$$\sigma_{\ell,\text{min}}^{\text{pt}} \propto \frac{\lambda}{d} \left(\frac{\sigma_t}{\tau_c^2 T_{\text{obs}}} \right)^{1/4}, \quad (56)$$

where $R \propto d^2$ has been used. Aperture scales as d^{-1} rather than d^{-2} , and observing time as $T_{\text{obs}}^{-1/4}$ rather than $T_{\text{obs}}^{-1/2}$ — both costs of the quadratic dependence on σ_ϕ . Each additional decade of sensitivity in σ_ℓ requires four decades of integration time, compared with two for the extended-source variant.

d. Continuum sources in the narrow-band limit. For a continuum source whose spectrum is flat within the bandpass, the coherence function is set entirely by the filter: a filter of width $\Delta \lambda$ gives $\tau_c = \lambda^2/(c \Delta \lambda)$, and the narrow-band Siegert discrepancy (Eq. 53) applies with this filter-defined τ_c . The Sun, observed through a single spatial coherence cell, illustrates the point-source case at extreme proximity. At $D = 1$ AU and $\lambda = 500$ nm, $\sigma_\phi(\alpha = 1/2) \approx 2 \times 10^{-5}$, giving $\mathcal{S}(\tau_c) \approx 1.1 \times 10^{-10}$. With a 1 Å filter ($\tau_c \approx 8$ ps, comparable to σ_t), the per-mode photon rate on a 10 m telescope is $r \sim 10^8$ cps — enormous compared to any extragalactic target. The single-channel source-limited SNR is $|\mathcal{S}| r \sqrt{\tau_c T_{\text{obs}}} \sim 2 \times 10^{-4}$ for $T_{\text{obs}} = 10^4$ hr. A spectrograph that disperses the light into N independent channels of width $\Delta \lambda$ measures $g^{(2)}$ in each channel independently and combines them for a $\sqrt{N}$ gain; with $N = 3000$ channels spanning the optical band, the combined SNR improves by $\sim 55 \times$ to $\sim 10^{-2}$ — still well below detection. The product $r \tau_c$ in each channel is fixed by the per-mode occupation number and is independent of the channel width, so no single-channel filter optimization helps; only the multiplexing gain $\sqrt{N}$ is available. The deficit is set entirely by the small $\sigma_\phi^2 \propto D^{2(1-\alpha)}$ at 1 AU, which is $\sim 10^{16}$ times smaller than at 1 Gpc.

2. Transverse Decorrelation – Zero-lag variant

For a spatially extended source whose angular size satisfies $D \Delta \theta \gtrsim r_\perp$ (equivalently $\Phi_{\text{eff}} \rightarrow 0$), photon pairs from different sky directions accumulate independent foam jitters and the zero-lag bunching peak is suppressed below its canonical value of 2:

$$g_{\text{obs}}^{(2)}(0) = 1 + e^{2s^2} \text{erfc}(\sqrt{2}s), \quad s \equiv \sigma_0 / \tau_L, \quad (57)$$

with $\sigma_0^2 = 2\sigma_\ell^2(D)(1 - \Phi_{\text{eff}})/c^2$ (Eqs. 47–45). The foam signal is the zero-lag Siegert discrepancy

$$\mathcal{S}(0) = g_{\text{obs}}^{(2)}(0) - 2 \leq 0, \quad (58)$$




FIG. 1. Signal-to-noise ratio for detecting the finite-lag Siegert discrepancy $\mathcal{S}(\tau_c) = e^{-2}(1 - e^{-2\sigma_\phi^2})$ (Eq. 36) as a function of foam exponent α , in the dynamical-foam limit $C_\alpha \rightarrow 0$. Instrument: Kim et al. [20] baseline (10 m aperture, $\epsilon = 0.39$, $\sigma_t = 13$ ps rms), $T_{\text{obs}} = 10$ hr, filter width $2 \times \text{FWHM}_{\text{line}}$. SNR computed from Eq. 54 (detector-limited regime, $\sigma_t \gg \tau_c$ for all sources shown). Filled circles mark the $\sigma_\phi = 0.3$ validity boundary of the weak-foam quadratic approximation $\mathcal{S}(\tau_c) \approx 2e^{-2}\sigma_\phi^2$; curves to the right saturate at $e^{-2} \approx 0.135$. Vertical dashed lines indicate $\alpha = 0.54$, the Lieu–Hillman imaging exclusion boundary ($\sigma_\phi = 1$ at $D = 1$ Gpc, $\lambda = 500$ nm, under the standard short- $r_\perp$ assumption), and $\alpha = 2/3$ (holographic). The finite-lag SNR tracks $D^{-2\alpha}$ (Eq. 55); its absolute sensitivity is orders of magnitude lower than the zero-lag variant because $\sigma_\phi \propto \sigma_t/\lambda$ is 10^{12} times smaller than $\sigma_0/\tau_L \propto \sigma_t/c\tau_L$ for optical lines. Figure generated by `make_snr_alpha_figure.py`.

a deficit relative to the Siegert prediction $g_{\text{src}}^{(2)}(0) = 2$. In the weak-foam limit $s \ll 1$ relevant to all viable foam models, $|\mathcal{S}(0)| = \delta g^{(2)} \approx 2\sqrt{2/\pi} s$ scales *linearly* in s — a consequence of the cusp in $g_{\text{src}}^{(2)}$ at $\tau = 0$ for sources with a Lorentzian line component (Eq. 48, Section VIE). The signal saturates at $|\mathcal{S}(0)| \rightarrow 1$ for $s \gg 1$.

Two structural conditions distinguish this variant from the finite-lag case and govern its applicability.

a. Source-size requirement. The mechanism requires $\Phi_{\text{eff}} < 1$, equivalently $D\Delta\theta \gtrsim r_\perp$: rays from different sky directions must sample distinct transverse regions of the foam to accumulate independent jitters. For sources unresolved on the foam transverse scale ($\Phi_{\text{eff}} \rightarrow 1$), $\sigma_0 \rightarrow 0$ and the variant becomes inactive; the finite-lag variant of Section VII A 1 then provides the only HBT channel. Because $r_\perp$ is itself an unknown parameter of the foam model, the source-size dependence of $\delta g^{(2)}$ across a sample of targets with varying $D\Delta\theta$ is the principal observational handle on $r_\perp$ (Section VI).

b. Filter strategy: preserving the intrinsic cusp. The linear-in- s scaling of $\delta g^{(2)}$ requires that the filter pass the full emission line,

$$\Delta\lambda_{\text{filter}} \geq \text{FWHM}_{\text{line}}, \quad (59)$$

so that the post-filter coherence time is set by the intrinsic Lorentzian line width rather than the (smooth) filter profile; a narrower filter destroys the cusp in $g_{\text{src}}^{(2)}$ at zero lag and reverts $\delta g^{(2)}$ to $O(s^2)$.

c. Detector-limited SNR. With $\sigma_t \gg \tau_L$ for all sources of interest, the bunching peak of width $\sim \tau_L$ is unresolved: the measured zero-lag excess in a time bin of width σ_t is diluted to $|\mathcal{S}(0)|_{\text{meas}} = (\tau_L/\sigma_t)|\mathcal{S}(0)|$, with statistical noise

$\sigma[g_{\text{meas}}^{(2)}] = 1/(R\sqrt{\sigma_t T_{\text{obs}}})$ [21–23], so

$$\text{SNR}_{\text{det}}^{\text{ext}} = \frac{\tau_L}{\sigma_t} |\mathcal{S}(0)| \cdot R\sqrt{\sigma_t T_{\text{obs}}} \xrightarrow{s \ll 1} 2\sqrt{\frac{2}{\pi}} \sigma_0 R\sqrt{T_{\text{obs}}/\sigma_t}. \quad (60)$$

The SNR in this regime is independent of τ_L : the cusp signal $\propto 1/\tau_L$ is cancelled by the τ_L/σ_t dilution factor. The penalty relative to source-limited operation ($\text{SNR}_{\text{src}}^{\text{ext}} = |\mathcal{S}(0)| R\sqrt{\tau_L T_{\text{obs}}}$, attainable only when $\sigma_t \lesssim \tau_L$) is $\sqrt{\tau_L/\sigma_t}$.

d. Distance and aperture scaling. Substituting $R \propto D^{-2}$ at fixed luminosity and $\sigma_0 \propto \sigma_\ell(D) \propto D^{1-\alpha}$ into Eq. 60 gives, in the weak-foam limit,

$$\text{SNR}^{\text{ext}} \propto D^{1-\alpha} \cdot D^{-2} = D^{-(1+\alpha)}, \quad (61)$$

with the penalty steeper for larger α ($D^{-3/2}$ for $\alpha = 1/2$, $D^{-5/3}$ for $\alpha = 2/3$). In the strong-foam limit $|\mathcal{S}(0)| \rightarrow 1$ saturates and $\text{SNR} \propto D^{-2}$. A nearby source of modest luminosity but high apparent flux can therefore be competitive with a more luminous but distant source: the higher photon rate R at close range can compensate for the reduced foam signal $|\mathcal{S}(0)|$.

The minimum detectable path-length amplitude follows from inverting Eq. 60 with the linear signal dependence $|\mathcal{S}(0)| \propto \sigma_0 \propto \sigma_\ell$:

$$\sigma_{\ell, \text{min}}^{\text{ext}} \propto \frac{\sigma_t}{d^2 \sqrt{\tau_L T_{\text{obs}}}}, \quad (62)$$

where $R \propto d^2$ has been used: sensitivity improves linearly with σ_t in the detector-limited regime and scales as d^{-2} (collecting area) and as $T_{\text{obs}}^{-1/2}$, both more favourable than the corresponding d^{-1} and $T_{\text{obs}}^{-1/4}$ scalings of the finite-lag variant. Importantly, this variant is independent of the foam temporal correlation C_α : the cross-ray differential jitter that drives $\delta g^{(2)}(0)$ operates regardless of whether the foam is dynamically active or quasi-static, in contrast to the finite-lag variant which requires $\tau_{\text{foam}} \lesssim \tau_c$.

e. Current- and next-generation sensitivity. With current instrumentation ($\sigma_t = 13$ ps), the broad-line quasar 3C 273 ($\tau_L \approx 0.11$ ps) operates deeply in the detector-limited regime. From Eq. 60 at 3σ ($R \approx 1.4 \times 10^5$ cps, $T_{\text{obs}} = 10^4$ hr): $\sigma_{\ell, \text{min}} \approx 1.7 \mu\text{m}$. This is far above the Lieu–Hillman imaging floor ($\sigma_\ell \approx 105$ nm at $\text{H}\alpha$), reflecting the intrinsic weakness of the envelope-timing observable $\delta g^{(2)}(0)$ relative to the carrier-phase observable $\sigma_\phi = 2\pi\sigma_\ell/\lambda$ that governs fringe visibility.

The detector-limited penalty $\sqrt{\tau_L/\sigma_t}$ vanishes only when σ_t is reduced below τ_L , which for broad-line quasars demands sub-0.1 ps timing. Superconducting nanowire detectors achieving such jitter are under active development; at $\sigma_t = 0.05$ ps, broad-line quasars enter the source-limited regime and NLS1s ($\tau_L \sim 0.5$ ps) are deeply source-limited, giving $\sigma_{\ell, \text{min}} \approx 160$ nm (10 m), 10 nm (40 m), 0.02 nm (1 km array) on 3C 273. A 40 m-class telescope with sub-ps timing would reach well below the Lieu–Hillman floor via $\delta g^{(2)}(0)$ alone, and a kilometre-baseline intensity-interferometer array would achieve sub-nanometre sensitivity — placing the random-walk model $\alpha = 1/2$ with long-range transverse correlations [11] firmly within reach across the full AGN/quasar redshift range.

f. Comparison of the two variants. For the 3C 273 reference parameters ($R \approx 1.4 \times 10^5$ cps, $\tau_c \approx 0.11$ ps, $T_{\text{obs}} = 10^4$ hr, $\lambda = \text{H}\alpha$, $D = 677$ Mpc), the random-walk model $\alpha = 1/2$ with short-range transverse correlations relaxed gives $\sigma_\phi \approx 230$, saturating the finite-lag signal at $\mathcal{S}(\tau_c) \rightarrow e^{-2}$ and $\sigma_{\ell, \text{min}}^{\text{pt}} \sim 10^2$ nm — comparable to the Lieu–Hillman fringe-erasure threshold, consistent with the finite-lag observable accessing σ_ϕ directly. For the holographic model $\alpha = 2/3$, $\sigma_\phi \sim 10^{-8}$ gives $\mathcal{S}(\tau_c) \sim 10^{-16}$, five orders of magnitude below the zero-lag signal $\delta g^{(2)}(0) \sim 10^{-11}$ on the same target. The finite-lag variant is therefore not competitive with the zero-lag variant for raw sensitivity to the holographic model, but it provides the only HBT channel sensitive to the foam temporal dynamics τ_{foam} and can be applied to spatially unresolved sources for which the zero-lag mechanism is inactive.

B. Candidate Sources

Three source properties drive target selection for the extended-source observable of Eq. (60): the photon rate R , set by the total line flux captured within the wide filter; the source distance, which sets σ_0 and hence $|\mathcal{S}(0)| = \delta g^{(2)} \approx 2\sqrt{2/\pi} \sigma_0/\tau_L$; and the line profile, which must retain a Lorentzian core to preserve the cusp in $g_{\text{src}}^{(2)}$ at zero lag and hence the linear-in- s scaling (Section VI E). Narrower lines (shorter τ_L) yield a larger cusp parameter $s = \sigma_0/\tau_L$ per unit foam amplitude and are therefore preferred targets at fixed photon rate. All sources discussed below have intrinsic coherence times $\tau_L \lesssim 1$ ps, well below the baseline timing resolution $\sigma_t = 13$ ps; the measurement therefore operates in the detector-limited regime throughout. The source classes surveyed below are illustrative rather than exhaustive; systematic searches through existing AGN flux catalogues and emission-line surveys will likely identify




FIG. 2. Signal-to-noise ratio for detecting the extended-source zero-lag bunching suppression $\delta g^{(2)}(0)$ as a function of foam exponent α , for five representative targets and $T_{\text{obs}} = 10$ hr on the Kim et al. [20] baseline instrument (10 m aperture, $\epsilon = 0.39$, $\sigma_t = 13$ ps rms). The foam signal is $|\mathcal{S}(0)| = 1 - e^{2s^2} \text{erfc}(\sqrt{2}s)$ with $s = \sigma_0/\tau_L$ (Eq. 48), and the SNR is computed from Eq. 60 (detector-limited regime, $\sigma_t \gg \tau_L$ for all sources shown). Photon rates use a wide filter of width $2 \times \text{FWHM}_{\text{line}}$, capturing $(2/\pi) \arctan(2) \approx 70\%$ of the Lorentzian line flux. Filled circles mark the $s = 0.3$ validity boundary of the weak-foam linear approximation $|\mathcal{S}(0)| \approx 2\sqrt{2/\pi}s$; curves to the left of the marker are in the linear regime. Vertical dashed lines indicate $\alpha = 0.54$ and $\alpha = 2/3$ (holographic). Figure generated by `make_snr_alpha_figure.py`.

superior targets with higher photon rates, more purely Lorentzian line profiles, or more favorable distance-to-flux ratios.

Broad-line quasars. Luminous broad-line quasars (BLQs) combine large σ_0 (from cosmological distances) with high line flux. Their permitted emission lines ($\text{H}\alpha$, $\text{H}\beta$, Mg II , C IV) concentrate photons into a narrow wavelength range; a wide filter ($\Delta\lambda_{\text{filter}} \geq \text{FWHM}$) captures the full line flux, giving photon rates far exceeding those of a continuum source at the same apparent magnitude. Sky background is negligible because the moderate filter width and small field of view suppress continuum sky to well below the source rate; detector dark counts ($\lesssim 25$ cps for premium SPADs, $\lesssim 1$ cps for SNSPDs) are well below typical bright AGN line rates, so the Brouri et al. correction for accidental coincidences is small [24]. We adopt 3C 273 (broad $\text{H}\alpha$, $F \approx 2 \times 10^{-12} \text{ erg s}^{-1} \text{ cm}^{-2}$, $\text{FWHM} \approx 127 \text{ \AA}$, $D = 677 \text{ Mpc}$) as the reference target [25]. Broad permitted lines in luminous quasars are empirically well described by Lorentzian profiles [26]. With a $\sim 200 \text{ \AA}$ filter (capturing $\approx 70\%$ of a Lorentzian of $\text{FWHM} = 127 \text{ \AA}$), the photon rate on a 10 m telescope with $\epsilon = 0.39$ is $R \approx 1.4 \times 10^5$ cps. The intrinsic coherence time is $\tau_L = \lambda^2/(c \text{FWHM}) = (6563 \text{ \AA})^2/(3 \times 10^{10} \text{ cm s}^{-1} \cdot 127 \text{ \AA}) \approx 0.11$ ps, placing the measurement deeply in the detector-limited regime ($\tau_L \ll \sigma_t = 13$ ps) with current instrumentation.

Narrow-line Seyfert 1 galaxies. Narrow-line Seyfert 1 galaxies (NLS1s) have permitted lines with $\text{FWHM} \lesssim 2000 \text{ km s}^{-1}$, empirically well described by a Lorentzian core [26]. A narrower line yields a shorter intrinsic coherence time τ_L and hence a larger cusp parameter $s = \sigma_0/\tau_L$ per unit foam amplitude: for $\text{FWHM} \sim 1300 \text{ km s}^{-1}$ ($\sim 28 \text{ \AA}$ at $\text{H}\alpha$), $\tau_L \approx 0.51$ ps, compared with $\tau_L \approx 0.11$ ps for 3C 273. For representative NLS1s with $\text{H}\alpha$ line fluxes $F \sim 5 \times 10^{-13} - 10^{-12} \text{ erg s}^{-1} \text{ cm}^{-2}$ at $D \sim 50 - 260 \text{ Mpc}$, photon rates with a wide filter on a 10 m telescope are $R \sim 10^4 - 10^5$ cps. Because NLS1 distances are smaller, their σ_0 is correspondingly smaller, but the genuinely

Lorentzian core and the high photon rates make them competitive targets for the $\delta g^{(2)}$ measurement: the Lorentzian purity maximises the linear-in- s cusp advantage of Section VIE, and NLS1s with $\tau_L \gtrsim \sigma_t$ can operate in the source-limited regime with next-generation detectors. Bright examples include Mrk 766 ($V = 13.6$, $D = 53$ Mpc, FWHM ≈ 1300 km s $^{-1}$) and I Zw 1 ($V = 14.0$, $D = 261$ Mpc, FWHM ≈ 1240 km s $^{-1}$).

Galactic emission-line sources as calibrators. Galactic sources — Wolf-Rayet stars with electron-scattering-broadened emission lines (e.g. γ^2 Vel, $D \approx 340$ pc; WR 6, $D \approx 1.7$ kpc), luminous blue variables (e.g. η Car, P Cyg), and symbiotic binaries with Stark-broadened H α (e.g. R Aqr, CH Cyg, $D \lesssim 1$ kpc) — can provide high photon rates and in several cases genuinely Lorentzian line cores from pressure or electron-scattering broadening. However, even for the most optimistic foam model ($\alpha = 1/2$), these sources have $\sigma_\phi \lesssim 0.3$ at kiloparsec distances: the foam signal is far below saturation and the predicted $\delta g^{(2)}$ is negligible. Their role is therefore as *systematic calibrators*: a non-detection on a bright nearby source with a well-characterised Lorentzian line profile confirms the absence of instrumental artefacts in the $g^{(2)}(0)$ measurement.

Sensitivity scalings $\sigma_{\ell, \min}^{\text{ext}}$ and the current/next-generation numerical sensitivities for 3C 273 and NLS1-class targets are given in Section VII A 2 (Eqs. 60–62); we do not repeat them here.

C. AGN/Quasar Redshift Sample as a Distance-Scaling Discriminant

Observing broad-line AGN and quasars across a wide redshift baseline provides a direct test of the distance scaling $\sigma_0 \propto \sigma_\ell(D) \propto D^{1-\alpha}$ predicted by the α -foam model. For representative targets spanning $z \approx 0.003$ to $z \approx 1.1$ (propagation distances $D \sim 14$ Mpc to ~ 4.7 Gpc, a factor of ~ 340 in D), the predicted $\delta g^{(2)}$ signal varies by more than an order of magnitude across the sample for $\alpha \approx 0.54$. The ratio $\delta g^{(2)}(D_{\text{far}})/\delta g^{(2)}(D_{\text{near}})$ directly encodes the D dependence of σ_0 and separates the foam exponent α from the coefficient A_α . A localized scatterer (thin-screen model) predicts $\sigma_\ell \propto D^{1/2}$, distinct from volumetric foam accumulation $\sigma_\ell \propto D^{1-\alpha}$ for $\alpha > 1/2$ (the two scalings coincide at the random-walk endpoint $\alpha = 1/2$); this distance-scaling test is therefore a model discriminant that operates independently of the absolute normalization A_α . The nearby Seyfert 1 galaxies (NGC 4151, Mrk 509) serve as systematic-calibration targets where the predicted foam signal is negligible, allowing any residual instrumental artefact in $g^{(2)}(0)$ to be characterized and subtracted.

D. Galactic and Local-Group Transient Targets

The cosmological-baseline strategy of Section VII C relies on sustained 10^4 hr campaigns on persistent broad-line AGN. A complementary strategy exploits rare local transients with strong narrow Balmer emission — specifically Type IIIn supernovae and luminous-blue-variable (LBV) giant eruptions — where the photon-flux advantage at $\sim$ kpc–Mpc distances dramatically exceeds that of any cosmological-distance source, even after derating for source luminosity, duty cycle, and event probability.

For a standard candle of fixed luminosity, $\text{SNR} \propto D^{-(1+\alpha)}$ in the weak-foam limit (Section VII A 2). For $\alpha = 2/3$ this is $\propto D^{-5/3}$: closer is dramatically better. Type Ia supernovae — the brightest cosmological-distance standard candles — are nevertheless poor targets at peak because their spectra are continuum-dominated with no narrow emission line: no Lorentzian cusp is available to drive the linear-in- s signal. Type IIIn supernovae and LBV giant eruptions, by contrast, exhibit strong narrow H α (and He) emission throughout their bright phase from circumstellar-material interaction, with intrinsic coherence times $\tau_L \sim 0.3$ –1 ps (e.g. $\tau_L \approx 0.26$ ps for a 55 Å FWHM SN IIIn line, $\tau_L \approx 0.95$ ps for the 15 Å FWHM core of η Car) that, with next-generation detectors ($\sigma_t \lesssim 0.1$ ps, i.e. $\sigma_t \lesssim \tau_L$), would enter the source-limited regime; the baseline instrument with $\sigma_t = 13$ ps remains firmly detector-limited even for these narrower lines.

For a galactic transient at $D \sim 2$ kpc with $\alpha = 1/2$ and long transverse correlations, the predicted $\sigma_\ell \approx 30$ nm and the photon flux is so high that the first-detection or exclusion is essentially instantaneous once sub-ps timing is available. For $\alpha = 2/3$, the $D^{2/3}$ signal penalty is offset by the photon-flux gain, with a galactic SN IIIn or LBV eruption returning a 10^3 – $10^4 \times$ SNR improvement over the persistent-AGN baseline.

The galactic SN IIIn case yields the largest single-event SNR — $\sim 10^4$ times the 3C 273 baseline — but its $\sim$ millennial recurrence makes it a “drop-everything” opportunity rather than a campaign target; the most recent confirmed Galactic core-collapse supernova was Cassiopeia A (~ 1680 , Type IIb). The η Car case sacrifices a factor of ~ 10 in SNR but offers a recurrence plausibly within a working career: η Car itself has been in an unusual brightening state since the mid-twentieth century, and giant LBV eruptions are now routinely identified in nearby galaxies (SN 2009ip and analogs). For both galactic sources, the $\sigma_\ell \approx 30$ nm prediction at $\alpha = 1/2$ with long-range transverse correlations is within reach of the HBT framework: a near-instant test of the long- $r_\perp$ random-walk variant. The M31 case is

intermediate: the $\alpha = 1/2$ short- $r_\perp$ prediction ($\sigma_\ell \approx 625$ nm) is already incompatible with existing imaging of M31, but the long- $r_\perp$ variant remains viable, and the $\alpha = 2/3$ photon-flux advantage is real if modest.

Two caveats apply specifically to η Car. The Homunculus Nebula scatters substantial flux into the line of sight, so the observed coherence is partially set by scattering geometry rather than the central engine; this must be modeled to credit the photon count correctly. During an eruption the photospheric envelope reaches ~ 2 mas at 2.3 kpc within ~ 10 yr at $v \sim 500$ km s $^{-1}$, resolved by intensity-interferometry baselines $\gtrsim 50$ m; the source is therefore no longer a clean point for the Section V Siegert test but does cleanly support the Section VI extended-source $\delta g^{(2)}$ test, with a directly imageable angular structure to anchor Φ_{eff} .

A standing rapid-response HBT plan for η Car, for any galactic core-collapse supernova, and for any nearby ($D \lesssim 30$ Mpc) Type II_n discovered by transient surveys would let the framework take advantage of these high-yield events without prior allocation.

E. Absorption-Line Sources

The framework applies without modification to sources observed in absorption. Absorption is a linear spectral filter applied to the continuum: the transmitted field amplitude is multiplied by a frequency-dependent real factor, which preserves the Gaussian chaotic character of the light. The transmitted power spectrum is $S_{\text{obs}}(\nu) = S_{\text{cont}}(\nu)T(\nu)$, where $T(\nu) = e^{-\tau(\nu)}$ is the transmission (for weak absorption, $T \approx 1 - \tau$, but the Sirius B example below has $\tau_{\text{central}} \approx 1.2$). Via the Wiener-Khinchin theorem, $g^{(1)}(\tau)$ is the Fourier transform of S_{obs} , and the Siegert relation and its foam violation apply exactly as before (Eq. 36). For $\mathcal{S}(\tau_c)$ the relevant observable is $\sigma_\phi(D)$, which depends only on the source distance and foam model, not on whether the source is observed in emission or absorption. Unlike a narrow emission line, which concentrates photons and enhances R by the inverse line-fraction factor, a broad absorption line can only reduce R relative to the unobscured continuum: there is no photon-rate advantage from the absorption profile.

The practical consequence is illustrated by the best-case absorption-line target. Sirius B (DA white dwarf, $V = 8.44$, $D = 2.64$ pc) has a deep, Stark-broadened H α profile (FWHM ~ 20 Å, central depth $\sim 70\%$). On a 10 m telescope with a 1.1 Å filter, the transmitted continuum rate is $R_{\text{cont}} \approx 8.8 \times 10^5$ cps, reduced to $R_{\text{abs}} \approx 2.6 \times 10^5$ cps at line center. At $D = 2.64$ pc and $\alpha = 1/2$, $\sigma_\phi \approx 0.011 \ll 1$, giving $\mathcal{S}(\tau_c) \approx 3.3 \times 10^{-5}$ and $\text{SNR}/\sqrt{T}/\text{hr} \approx 1.9 \times 10^{-3}$ — about $17\times$ smaller than a 3C 273-class AGN in the saturated regime ($\mathcal{S} \rightarrow e^{-2}$, $\text{SNR}/\sqrt{T}/\text{hr} \approx 3.3 \times 10^{-2}$). For $\alpha = 2/3$, $\sigma_\phi \approx 2.7 \times 10^{-11}$ at 2.64 pc: the foam signal is entirely negligible. Absorption-line sources are therefore not competitive as foam probes; their value lies in null-test calibration, where a nearby DA white dwarf provides a Lorentzian-profile source with negligible predicted foam signal for any plausible α .

The solar Ca II K line ($\lambda = 3934$ Å) illustrates the combined limits of proximity and extended-source dilution. The photospheric damping wings are pressure-broadened with a Lorentzian FWHM $\Gamma_{\text{abs}} \approx 0.3$ Å, but the Lorentzian wing absorption at $\Delta\lambda \sim 2$ Å from line center is only $\sim 0.6\%$; filters placed in the damping wing are observing essentially pure continuum with negligible line-profile modulation. The Sun is also an extended source: a 10 m aperture collects $M = A_{\text{tel}}/A_{\text{coh}} \approx 4 \times 10^{10}$ independent spatial modes (coherence area $A_{\text{coh}} = (\lambda/\theta_\odot)^2 \approx 1.8 \times 10^{-9}$ m 2 for angular diameter $\theta_\odot \approx 9.3$ mrad), each contributing an independent foam phase. The SNR formula (Eq. 60) therefore applies with the per-mode rate $r = R_{\text{total}}/M = 2\bar{n}(\lambda)\Delta\nu\epsilon$ (Section VI), where $\bar{n}(\lambda_K) = [e^{hc/\lambda_K k_B T} - 1]^{-1} \approx 1.8 \times 10^{-3}$ for $T = 5778$ K, giving $r \approx 1.1 \times 10^8$ cps ($\Delta\nu = \sigma_t^{-1}$, two polarisations). Despite this, the foam signal is minuscule: $\sigma_\phi(1 \text{ AU}, \lambda_K, \alpha = 1/2) \approx 2.5 \times 10^{-5}$, giving $\mathcal{S}(\tau_c) \approx 1.7 \times 10^{-10}$ and $\text{SNR}/\sqrt{T}/\text{hr} \approx 2 \times 10^{-6}$, about $10^4\times$ below the AGN reference. The root cause is proximity: $\sigma_\phi \propto D^{1/2-\alpha}$ at $D = 1$ AU is $\sim 10^8$ times smaller than at $D = 1$ Gpc, and no photon-rate advantage can compensate.

VIII. CONCLUSIONS

We have developed a theoretical framework for the imprint of spacetime foam, characterized by the foam exponent α , on both first- and second-order optical coherence. The central finding is that foam-induced arrival-time jitter produces two classes of observable violation of relations that hold for all Gaussian chaotic sources: a deviation from the Siegert relation, $g^{(2)}(\tau) - 1 = |g^{(1)}(\tau)|^2$, which requires the foam to be dynamically active (τ_{foam} short relative to the optical coherence time), and a suppression of the zero-lag bunching peak below its canonical value of $g^{(2)}(0) = 2$, which depends only on $\sigma_\ell^2(D)$ and is independent of τ_{foam} . These second-order observables are complementary to the fringe visibility and Strehl ratio tests based on first-order coherence that have already been applied to constrain foam models. For the holographic model ($\alpha = 2/3$), all such observables lie well within the currently inaccessible parameter space: the first-order tests probe effects at the $\sim 10^{-8}$ level, while the second-order observables are characteristically

smaller at $\sim 10^{-11}$ (a recent analysis of GRB 221009A claims consistency with holographic foam blurring [15], though without formal statistical significance). For the holographic model the predicted signals at $\sim 10^{-11}$ lie well below the sensitivity of current and near-term instrumentation; foam models with $1/2 \ll \alpha < 2/3$ predict larger signals and may be within reach.

The advantage of the second-order approach lies not in raw sensitivity but in two distinct features. First, the Siegert discrepancy $\mathcal{S}(\tau) = g^{(2)}(\tau) - 1 - |g^{(1)}(\tau)|^2$ is constructed directly from the data; a detection of $\mathcal{S} \neq 0$ can use the measured source spectrum to obtain $g_{\text{src}}^{(1)}(\tau)$ via the Wiener–Khinchin theorem, without a separate source model. Second, for extended sources whose line profile has a Lorentzian component (natural or pressure broadening), the zero-lag bunching suppression $\delta g^{(2)}$ scales linearly with $s = \sigma_0/\tau_L$ — a consequence of the cusp in $g_{\text{src}}^{(2)}$ at zero lag — so that despite being a second-order coherence measurement, $\delta g^{(2)}$ is not doubly suppressed relative to the first-order observables. Physically, the $\delta g^{(2)}$ suppression reflects the two-photon character of HBT bunching: the zero-lag peak is built from photon pairs arriving from different sky directions $\theta \neq \theta'$, which propagate through distinct transverse regions of the foam and accumulate independent jitters $\delta t(t; \theta) \neq \delta t(t; \theta')$ when $b \gg r_\perp$; their relative offset blurs $g^{(2)}(0)$ below its canonical value of 2 (cf. Section VI E).

For the fiducial parameters $D = 1$ Gpc, $\lambda = 500$ nm, $\Delta\lambda/\lambda = 10^{-3}$, $A_\alpha = 1$: the $\alpha = 1/2$ case gives $\sigma_\phi \approx 280$ (fringes completely erased) and is excluded by fringe-visibility measurements [8–10] under the standard assumption of short-range transverse foam correlations; long-range-correlated random-walk foam [11] evades this bound and remains a target of the present framework. The holographic case $\alpha = 2/3$ gives $\sigma_\phi \sim 10^{-8}$ — six orders of magnitude below the $\sigma_\phi \sim 1$ imaging exclusion — predicting $\mathcal{S}/|g_{\text{src}}^{(1)}|^2 \sim 10^{-15}$ for the finite-lag Siegert violation and $\delta g^{(2)} \approx 2\sqrt{2/\pi}s \sim 5.7 \times 10^{-11}$ for the extended-source bunching suppression in the fully resolved limit $\Phi_{\text{eff}} \rightarrow 0$, where $s = \sqrt{2}\sigma_\phi(\Delta\lambda/\lambda) \sim 3.5 \times 10^{-11}$ (with $\delta g^{(2)}$ inheriting a linear rather than quadratic scaling in $s = \sigma_0/\tau_c$ from the cusp in $g_{\text{src}}^{(2)}$ at zero lag for sources with a Lorentzian line component).

The achievable sensitivity and its scaling with telescope aperture and source redshift are analyzed in detail in Section VII. The observable is the zero-lag bunching suppression $\delta g^{(2)}(0)$ (Eq. (48), $\Phi_{\text{eff}} \approx 0$), measured via the HBT split-beam configuration using bright broad-line AGN and quasars observed through wide filters that preserve the intrinsic Lorentzian line cusp. The linear-in- s scaling of $\delta g^{(2)}$ requires that the filter pass the full emission line ($\Delta\lambda_{\text{filter}} \geq \text{FWHM}_{\text{line}}$); this places bright quasars ($\tau_L \sim 0.1$ ps) in the detector-limited regime with current timing resolution ($\sigma_t = 13$ ps), with a $\sqrt{\tau_L/\sigma_t} \approx 0.09$ penalty relative to source-limited operation. For a 10-m-class telescope [20] with $T_{\text{obs}} = 10^4$ hr on a 3C 273-class quasar, the current-detector sensitivity is $\sigma_{\ell, \text{min}} \approx 1.7 \mu\text{m}$ (3σ) — above the Lieu–Hillman imaging floor ($\sigma_\ell \approx 105$ nm at H α). With next-generation timing ($\sigma_t \lesssim 0.1$ ps, entering the source-limited regime), a 40 m aperture reaches $\sigma_\ell \approx 10$ nm, well below the imaging floor; a 1 km-class intensity-interferometer array would reach sub-nm sensitivity.

The astrophysical contribution of the framework developed here is threefold: it closes the imaging-test loophole for foam models with long-range transverse correlations, it provides a wavelength-dependence discriminant against chromatic plasma scattering (the Siegert violation inherits the $\sigma_\phi \propto 1/\lambda$ scaling), and it enables $r_\perp$ inversion via $\delta g^{(2)}$ on nearby AGN whose broad emission lines have Lorentzian components. What HBT-on-AGN cannot do is reach the holographic prediction $\sigma_\phi \sim 10^{-8}$, which lies six orders of magnitude below the most aggressive HBT projection; closing this gap would require photon flux $\sim 10^{20}$ greater than from a luminous quasar. The holographic regime is the natural target of tabletop holographic-noise interferometers [27] that bypass the photon-counting penalty.

The $g^{(2)}(0)$ observable developed here is already routinely computed in radio astronomy under different names: spectral kurtosis [28, 29], used for real-time RFI excision at the VLA, GBT, LOFAR, and CHIME, is constructed to estimate $g^{(2)}(0) - 1$ from channelized power samples (returning unity in expectation for a stationary Gaussian field, up to finite-sample bias corrections), as does the modulation index $m^2 \equiv \text{Var}(I)/\langle I \rangle^2$, which equals $g^{(2)}(0) - 1$ for stationary, noise-corrected intensity samples and is used to characterize single-pulse intensity statistics of pulsars and FRBs [30]. The Siegert-violation framework developed here therefore translates directly to existing radio data products.

A deeper question left open by the present framework is whether the foam exponent α or coefficient A_α carries a redshift dependence. If so, observations of sources at different redshifts would deviate from the simple power-law $\sigma_\phi(D)$ predicted here, and the distance-scaling discriminant of Section VI would need to be generalized; such a deviation would itself be a cosmological signal. We flag this as motivation for a redshift-resolved treatment of foam coherence in future work.

ACKNOWLEDGMENTS

This work was supported in part by the Director, Office of Science, Office of High Energy Physics of the U.S. Department of Energy under Contract No. DE-AC02-05CH11231.

TABLE II. Observed properties of the AGN/quasar candidate sources. z is the spectroscopic redshift; D is the luminosity distance for $H_0 = 70 \text{ km s}^{-1} \text{ Mpc}^{-1}$; m_V is the apparent visual magnitude; line fluxes F are integrated over the emission line; FWHM is the observed-frame full width at half maximum of the broad emission line component; $\tau_c = \lambda_{\text{obs}}^2 / (c \Delta \lambda_{\text{FWHM}})$ is the coherence time of the unfiltered line.

Source	z	D (Mpc)	m_V	Line	F ($\text{erg s}^{-1} \text{ cm}^{-2}$)	FWHM ($\text{\AA}$)	τ_c (ps)
3C 273	0.158	677	12.9	H α	2×10^{-12} [25]	127 [25]	0.15
Mrk 766	0.013	53	13.6	H α	...	28 [26]	0.53
I Zw 1	0.061	261	14.0	H α	...	27 [26]	0.60
NGC 4151	0.003	14	11.5	H α	5×10^{-11}	...	...
Mrk 509	0.034	146	13.1	H α	...	...	...

The authors used Claude (Anthropic) to assist with drafting and editing portions of this manuscript. All content was reviewed and verified by the authors, who take full responsibility for the work.

Appendix A: Voigt-Profile Convolution

Equations (18)–(19) follow from evaluating Eq. (14) with the Voigt source coherence $g_{\text{src}}^{(1)}(\tau) = e^{-|\tau|/\tau_L} e^{-\tau^2/2\tau_G^2}$:

$$g_{\text{obs}}^{(1)}(\tau) = \int_{-\infty}^{\infty} \mathcal{N}(\Delta; 0, \sigma_{\Delta}^2) e^{-|\tau-\Delta|/\tau_L} e^{-(\tau-\Delta)^2/2\tau_G^2} d\Delta. \quad (\text{A1})$$

The two Gaussian factors $\mathcal{N}(\Delta; 0, \sigma_{\Delta}^2)$ and $e^{-(\tau-\Delta)^2/2\tau_G^2}$ combine into a single Gaussian in Δ ; completing the square gives the effective mean and variance

$$\tilde{\sigma}_1^2 = \frac{\sigma_{\Delta}^2 \tau_G^2}{\sigma_{\Delta}^2 + \tau_G^2}, \quad \mu_1(\tau) = \frac{\tau_G^2 \tau}{\sigma_{\Delta}^2 + \tau_G^2}, \quad (\text{A2})$$

together with the multiplicative prefactor $(\tilde{\sigma}_1/\sigma_{\Delta}) e^{-\tau^2/2(\sigma_{\Delta}^2 + \tau_G^2)}$. The remaining Lorentzian–Gaussian convolution $\int \mathcal{N}(\Delta; \mu_1, \tilde{\sigma}_1^2) e^{-|\tau-\Delta|/\tau_L} d\Delta$ is split at $\Delta = \tau$; each half-line integral is of the standard form $\int_0^{\infty} e^{-x/\tau_L} e^{-(x-a)^2/2\tilde{\sigma}_1^2} dx \propto e^{\tilde{\sigma}_1^2/2\tau_L^2 \mp a/\tau_L} \text{erfc}((\tilde{\sigma}_1/\tau_L \mp a/\tilde{\sigma}_1)/\sqrt{2})$ after completing the square in the exponent, and the two halves combine into Eq. (19). The pure-Lorentzian closed form Eq. (16) follows in the limit $\tau_G \rightarrow \infty$, for which $\tilde{\sigma}_1 \rightarrow \sigma_{\Delta}$, $\mu_1 \rightarrow \tau$, and the prefactor reduces to unity.

The zero-lag Voigt result Eq. (51) for the second-order coherence is obtained by the same procedure applied to $g_{\text{src}}^{(2)}(\tau) - 1 = e^{-2|\tau|/\tau_L} e^{-\tau^2/\tau_G^2}$: the factor-of-two changes in the exponents of both the Lorentzian and Gaussian factors propagate through, yielding the zero-lag effective width $\tilde{\sigma}^2 = \sigma_0^2 \tau_G^2 / (2\sigma_0^2 + \tau_G^2)$ of Eq. (50).

Appendix B: Source Properties

Tables II and III collect the observed properties of the principal candidate sources used in Section VII.

Broad-line profile shapes

The profile shape of the broad emission line determines the fraction of line flux transmitted through a narrow filter and the functional form of $g_{\text{src}}^{(2)}(\tau)$ at zero lag (Section VI E).

Sulentic et al. [31] and Zamfir et al. [32] establish an empirical division of AGN into two populations based on FWHM(H β): *Population A* sources with FWHM $\lesssim 4000 \text{ km s}^{-1}$ have broad H β profiles that are well described by a single Lorentzian, while *Population B* sources with FWHM $\gtrsim 4000 \text{ km s}^{-1}$ require a double-Gaussian (a classical broad component plus a very broad, redshifted component) and are not well described by a single Lorentzian.

Mrk 766 has FWHM $\approx 1300 \text{ km s}^{-1}$, placing it firmly in Population A; its broad H α profile is well described by a Lorentzian [31, 32].

3C 273 has FWHM $\approx 6800 \text{ km s}^{-1}$ (rest-frame velocity equivalent of the $127 \text{ \AA}$ observed-frame H β width at $z = 0.158$ [25]), placing it in Population B. Its broad H β profile is not well described by a single Lorentzian and instead requires a double-Gaussian decomposition [31, 32].

TABLE III. Observed properties of η Carinae during the Great Eruption, from light-echo spectroscopy [33, 34]. $D = 2.3$ kpc [35]. FWHM values are for the broad H α emission component; the bulk outflow speed characterises the P Cygni absorption trough or the dominant spectral feature at each epoch. $\tau_c = \lambda^2/(c\Delta\lambda_{\text{FWHM}})$ at H α ($\lambda = 6563$ Å) assuming a Lorentzian profile. The eruption reached $m_V \approx -1$ at peak (1843–1858). Later phases also show very broad emission wings with velocities $\gtrsim 10,000$ km s $^{-1}$ [33].

Eruption epoch	Bulk outflow (km s $^{-1}$)	H α FWHM (km s $^{-1}$)	τ_c (ps)	Ref.
Pre-plateau (early)	~ 150	~ 150	4.38	[33]
Plateau	~ 400	~ 400	1.64	[34]
Plateau peak	~ 600	~ 600	1.09	[33]

-
- [1] J. A. Wheeler, *Physical Review* **97**, 511 (1955).
[2] S. W. Hawking, *Nuclear Physics B* **144**, 349 (1978).
[3] S. Carlip, *Rep. Prog. Phys.* **86**, 066001 (2023), arXiv:2209.14282 [gr-qc].
[4] Y. J. Ng and E. S. Perlman, *Universe* **8**, 382 (2022).
[5] Y. J. Ng, *Physical Review Letters* **86**, 2946 (2001).
[6] Y. J. Ng, *Modern Physics Letters A* **18**, 1073 (2003).
[7] R. Lieu and L. W. Hillman, *The Astrophysical Journal* **585**, L77 (2003).
[8] R. Ragazzoni, M. Turatto, and W. Gaessler, *ApJ* **587**, L1 (2003), arXiv:astro-ph/0303043 [astro-ph].
[9] W. A. Christiansen, Y. J. Ng, D. J. E. Floyd, and E. S. Perlman, *Physical Review D* **83**, 084003 (2011).
[10] E. S. Perlman, S. A. Rappaport, W. A. Christiansen, Y. J. Ng, J. DeVore, and D. Pooley, *The Astrophysical Journal* **805**, 10 (2015).
[11] V. S. H. Lee, K. M. Zurek, and Y. Chen, *Physical Review D* **109**, 084005 (2024), arXiv:2312.06757 [gr-qc].
[12] A. A. Abdo, M. Ackermann, M. Ajello, *et al.* (Fermi GBM/LAT Collaboration), *Nature* **462**, 331 (2009), arXiv:0908.1832 [astro-ph].
[13] F. Aharonian *et al.* (H.E.S.S. Collaboration), *Phys. Rev. Lett.* **101**, 170402 (2008), arXiv:0810.3475 [astro-ph].
[14] LHAASO Collaboration, *Physical Review Letters* **133**, 071501 (2024), arXiv:2402.06009.
[15] E. Steinbring and Y. J. Ng, arXiv e-prints (2026), arXiv:2603.18277.
[16] R. H. Brown and R. Q. Twiss, *Nature* **177**, 27 (1956).
[17] A. J. F. Siegert, *On the fluctuations in signals returned by many independently moving scatterers*, Tech. Rep. 465 (MIT Radiation Laboratory, 1943).
[18] A. U. Abeysekara *et al.*, *Nature Astronomy* **4**, 1164 (2020).
[19] A. Acharyya *et al.*, *The Astrophysical Journal* **966**, 28 (2024).
[20] A. G. Kim, P. E. Nugent, X. Chen, L. Wang, J. T. O’Brien, and Tardis Collaboration, *Phys. Rev. D* **111**, 083047 (2025), arXiv:2503.07725 [astro-ph.IM].
[21] C. Foellmi, *Astronomy & Astrophysics* **507**, 1719 (2009).
[22] D. Dravins, S. LeBohec, H. Jensen, and P. D. Nun ez, *Astroparticle Physics* **43**, 331 (2013).
[23] W. Gu erin, A. Dussaux, M. Fouch e, G. Labeyrie, R. Donnerstein, D. Nakar, A. Gontier, and R. Kaiser, *Monthly Notices of the Royal Astronomical Society* **472**, 4126 (2017).
[24] R. Brouri, A. Beveratos, J.-P. Poizat, and P. Grangier, *Optics Letters* **25**, 1294 (2000).
[25] T. A. Boroson and R. F. Green, *The Astrophysical Journal Supplement Series* **80**, 109 (1992).
[26] W. Kollatschny and M. Zetzl, *Astronomy & Astrophysics* **549**, A100 (2013).
[27] E. P. Verlinde and K. M. Zurek, *Physics Letters B* **822**, 136663 (2021), arXiv:1902.08207 [gr-qc].
[28] G. M. Nita, D. E. Gary, Z. Liu, G. J. Hurford, and S. M. White, *Publications of the Astronomical Society of the Pacific* **119**, 805 (2007).
[29] G. M. Nita and D. E. Gary, *Monthly Notices of the Royal Astronomical Society* **406**, L60 (2010).
[30] P. Weltevrede, R. T. Edwards, and B. W. Stappers, *Astronomy and Astrophysics* **445**, 243 (2006).
[31] J. W. Sulentic, P. Marziani, R. Zamanov, R. Bachev, M. Calvani, and D. Dultzin-Hacyan, *The Astrophysical Journal* **566**, L71 (2002).
[32] S. Zamfir, J. W. Sulentic, P. Marziani, and D. Dultzin, *Monthly Notices of the Royal Astronomical Society* **403**, 1759 (2010).
[33] N. Smith, J. E. Andrews, A. Rest, F. B. Bianco, J. L. Prieto, T. Matheson, D. J. James, R. C. Smith, G. M. Strampelli, and A. Zenteno, *Monthly Notices of the Royal Astronomical Society* **480**, 1466 (2018).
[34] J. L. Prieto, A. Rest, F. B. Bianco, T. Matheson, N. Smith, N. R. Walborn, E. Y. Hsiao, R. Chornock, L. Paredes  lvarez, A. Campillay, C. Contreras, C. Gonz lez, D. James, G. R. Knapp, A. Kunder, S. Margheim, N. Morrell, M. M. Phillips, R. C. Smith, D. L. Welch, and A. Zenteno, *The Astrophysical Journal* **787**, L8 (2014).
[35] D. J. Hillier, K. Davidson, K. Ishibashi, and T. Gull, *The Astrophysical Journal* **553**, 837 (2001).