First- and Second-Order Coherence Through Spacetime Foam

A. G. Kim* and P. E. Nugent
Lawrence Berkeley National Laboratory
Berkeley, California 94720, USA

L. Wang
Department of Physics and Astronomy 4242 TAMU
College Station, Texas 77843-4242, USA
(Dated: April 24, 2026)

We study first- and second-order coherence of light from point and extended sources propagating through spacetime foam characterized by the Ng–Perlman exponent α. In the narrow-band regime, arrival-time jitter suppresses temporal first-order coherence and produces a Siegert violation whose foam dependence encodes α and the line-of-sight path-length scale $\sigma_\ell(D)$. For extended sources, transverse foam coherence introduces an additional suppression of the zero-lag Hanbury Brown–Twiss bunching peak that directly probes the transverse coherence length $r_\perp$. We develop the theoretical framework through which joint measurements of the Siegert discrepancy and zero-lag bunching suppression across sources of different angular size and redshift could constrain $\{\alpha, \sigma_\ell(D), r_\perp\}$ simultaneously, and identify the distance scaling and achromaticity as discriminants against conventional propagation phase screens.

I. MODEL

A. Arrival-Time Jitter as a Time Reparameterization

We model the detected field as a stochastically time-warped version of the intrinsic source field:

$$E_{\rm obs}(t) = E_{\rm src}(t - \delta t(t))\,, \tag{1}$$

where $E_{\rm src}(t)$ is a stationary zero-mean complex Gaussian process (a single chaotic/thermal mode) and $\delta t(t)$ is the foam-induced arrival-time jitter, a stationary stochastic process statistically independent of $E_{\rm src}$. The jitter arises from cumulative path-length fluctuations along the line of sight of propagation distance D:

$$\delta t(t) = \frac{1}{c}\,\Delta\ell_{\rm tot}(t). \tag{2}$$

In the geometric-optics limit the path-length shift $\Delta\ell_{\rm tot}$ is a purely geometric quantity and hence frequency-independent (achromatic), in contrast to plasma dispersion which scales as λ^2. This achromaticity is a key discriminant between foam and conventional propagation effects (see Section IV F).

B. Ng–Perlman Variance and Temporal Covariance

Following Ng and collaborators [1–3] (see also Ref. [4] for a review of foam models and their observational constraints), the single-ray path-length variance obeys the power-law scaling

$$\sigma_\ell^2(D) = A_\alpha\, D^{2(1-\alpha)}\, \ell_{\rm P}^{2\alpha}, \tag{3}$$

where $\ell_{\rm P} = \sqrt{\hbar G/c^3} \simeq 1.616 \times 10^{-35}$ m is the Planck length, α parametrizes the foam model ($\alpha = 1$: conservative; $\alpha = 2/3$: holographic; $\alpha = 1/2$: random walk), and A_α is a dimensionless coefficient of order unity. The two-time covariance of the aggregate path-length fluctuation along a single ray takes the factorized form

$$\big\langle \Delta\ell_{\rm tot}(t)\,\Delta\ell_{\rm tot}(t+\tau)\big\rangle = \sigma_\ell^2(D)\, C_\alpha(\tau; D), \tag{4}$$

* agkim@lbl.gov

where $C_\alpha(\tau; D)$ satisfies $C_\alpha(0; D) = 1$ and decays to zero as $\tau \to \infty$. This product form — separating the total variance from the normalized temporal shape — is a structural assumption of the model, not a consequence of Eq. (3); in a general stationary process the variance and the normalized correlation function can carry independent D-dependence.

Because foam is a property of spacetime rather than of any preferred direction, the same dimensional logic that governs longitudinal path-length fluctuations applies equally to transverse baselines. For two rays separated by a transverse distance b, the Ng–Perlman scaling gives a characteristic transverse coherence length

$$r_\perp \sim A_\alpha^{1/2}\, b^{1-\alpha}\, \ell_{\rm P}^{\alpha}, \tag{5}$$

by dimensional analogy with Eq. (3). This analogy motivates the *existence* of a transverse scale but does not derive its value: in the longitudinal case D counts the number of cells traversed along the ray, while in the transverse case b is a separation between rays that share cells along the line of sight — the number of shared cells depends on both b and D in a way that is geometrically distinct from the single-ray accumulation. Accordingly, $r_\perp$ is treated as an independent free parameter in Section V, where it is constrained empirically through the angular-size dependence of the HBT bunching suppression. The two principal observables developed below — the Siegert discrepancy $\mathcal{S}(\tau)$ and the zero-lag bunching suppression $\delta g^{(2)}$ — therefore both descend from the single scaling relation Eq. (3), with Sections II–IV probing its longitudinal projection and Section V its transverse projection, all governed by the same parameter set $\{\alpha,\, \ell_{\rm P},\, A_\alpha\}$.

C. Differential Jitter and Its Distribution

Define the differential jitter between two detection times separated by lag τ:

$$\Delta(\tau) \equiv \delta t(t+\tau) - \delta t(t). \tag{6}$$

For a fixed propagation distance D, stationarity implies that the statistics of $\Delta(\tau)$ depend only on τ. Its variance follows directly from Eq. (4):

$$\begin{aligned} \sigma_\Delta^2(\tau; D) &\equiv \langle \Delta(\tau)^2 \rangle \\ &= \frac{1}{c^2}\big\langle [\Delta\ell_{\rm tot}(t+\tau) - \Delta\ell_{\rm tot}(t)]^2 \big\rangle \\ &= \frac{2\sigma_\ell^2(D)}{c^2}\big[1 - C_\alpha(\tau; D)\big]. \end{aligned} \tag{7}$$

Because $\Delta(\tau)$ is a sum of contributions from many weakly correlated foam cells, the central limit theorem implies it is Gaussian with zero mean. We therefore write its conditional probability density (at fixed D) as

$$p(\Delta; \tau, D) = \frac{1}{\sqrt{2\pi}\,\sigma_\Delta(\tau; D)} \exp\!\left(-\frac{\Delta^2}{2\sigma_\Delta^2(\tau; D)}\right). \tag{8}$$

Note the two limits: $\sigma_\Delta^2(0; D) = 0$ (no jitter between simultaneous arrivals) and $\sigma_\Delta^2(\infty; D) = 2\sigma_\ell^2(D)/c^2$ (maximal decorrelation at large lag).

For notational compactness in subsequent sections, we suppress the explicit D-dependence and write $C_\alpha(\tau) \equiv C_\alpha(\tau; D)$ and $\sigma_\Delta(\tau) \equiv \sigma_\Delta(\tau; D)$ (similarly for any other D-dependent quantity).

D. Physical Origin of $C_\alpha(\tau; D)$

Quantum foam models posit that each Planck-scale cell of spacetime along the line of sight fluctuates in time. Two photons traversing the same cell at coordinate times separated by τ experience path-length fluctuations that are correlated for $\tau \ll \tau_{\rm foam}$ but decorrelated for $\tau \gg \tau_{\rm foam}$, where $\tau_{\rm foam}$ is the characteristic temporal fluctuation timescale at each cell. Consequently, $C_\alpha(\tau; D)$ is the normalized *temporal* autocorrelation of the per-cell path-length fluctuation, integrated over the $\sim D/\ell_{\rm P}$ cells along the line of sight; $\sigma_\ell^2(D)$ encodes the accumulated variance. The timescale $\tau_{\rm foam}$ is an additional free parameter of the foam model, not fixed by the Ng–Perlman variance relation alone.

If $\tau_{\rm foam}$ is of order the Planck time, $\ell_{\rm P}/c \approx 5 \times 10^{-44}$ s, then $C_\alpha(\tau) \approx 0$ for all accessible lags and $\sigma_\Delta^2(\tau) \approx 2\sigma_\ell^2/c^2$ at all τ: the maximum differential jitter, and the most optimistic scenario for detection. In the opposite *quasi-static* limit $\tau_{\rm foam} \gg D/c$ (foam coherent across the entire line of sight), $C_\alpha(\tau) \approx 1$ for all τ of interest, $\sigma_\Delta^2 \to 0$, and

every temporal observable vanishes. In this limit a null measurement places a *lower bound* on $\tau_{\rm foam}$; however, this bound is meaningful only when combined with an independent constraint on σ_ϕ — for example from fringe-visibility measurements, since a null HBT result alone is equally consistent with the observationally allowed regime $\sigma_\phi \ll 1$ where the foam signal is undetectably small regardless of $\tau_{\rm foam}$.

The quantity $\xi_\parallel \equiv c\,\tau_{\rm foam}$ is the distance light travels in one foam coherence time; it provides a length-scale interpretation of $\tau_{\rm foam}$ but is fundamentally temporal in origin. For cosmological D, a complete treatment would require accounting for the redshift when converting proper-time lags at each cell to observed lags; this correction is beyond the scope of the present paper. Furthermore, we treat the foam's statistical properties as stationary over the observation duration T; any secular drift due to cosmological expansion is of order $H_0 T \lesssim 10^{-10}$ and thus strictly negligible.

The Gaussian approximation for $\Delta(\tau)$ requires the elementary foam cells contributing to to $\Delta\ell_{\rm tot}$ to be approximately independent. This holds when the temporal lag of interest satisfies $\tau \gg \tau_{\rm foam}$ (so cells decorrelate over one lag interval) and when the transverse baseline $b \gg r_\perp$; if $\tau_{\rm foam}$ is comparable to the coherence time τ_c, or if transverse correlations extend over b, the Gaussian approximation may acquire non-Gaussian tails.

II. FIRST-ORDER COHERENCE

A. General Result

The observed first-order coherence is

$$g^{(1)}_{\rm obs}(\tau) = \frac{\langle E^*_{\rm obs}(t)\, E_{\rm obs}(t+\tau)\rangle}{\langle |E_{\rm obs}(t)|^2\rangle}. \tag{9}$$

Substituting Eq. (1) and letting $u = t - \delta t(t)$, $v = t + \tau - \delta t(t+\tau)$, so that $v - u = \tau - \Delta(\tau)$, stationarity of $E_{\rm src}$ gives

$$\langle E^*_{\rm src}(u)\, E_{\rm src}(v)\rangle = \langle |E_{\rm src}|^2\rangle\, g^{(1)}_{\rm src}(\tau - \Delta(\tau)). \tag{10}$$

Since δt is independent of $E_{\rm src}$, averaging over the jitter yields

$$\boxed{g^{(1)}_{\rm obs}(\tau) = \big\langle g^{(1)}_{\rm src}(\tau - \Delta(\tau))\big\rangle_{\delta t} = \int_{-\infty}^{\infty} p(\Delta;\tau)\, g^{(1)}_{\rm src}(\tau - \Delta)\, d\Delta.} \tag{11}$$

The observed first-order coherence is the intrinsic coherence function convolved with the Gaussian differential-jitter kernel.

B. Closed Form for a Lorentzian Source

Astronomical emission lines generally have a Lorentzian line profile, which corresponds to an exponential coherence function via the Wiener-Khinchin theorem, so

$$g^{(1)}_{\rm src}(\tau) = e^{-|\tau|/\tau_c}, \tag{12}$$

where τ_c is the coherence time. Evaluating Eq. (11) by completing the square and splitting on the sign of $\tau - \Delta$ gives, for $\tau > 0$:

$$g^{(1)}_{\rm obs}(\tau) = e^{\sigma_\Delta^2/2\tau_c^2} \cdot \frac{1}{2}\left[e^{-\tau/\tau_c}\,{\rm erfc}\!\left(\frac{\sigma_\Delta/\tau_c - \tau/\sigma_\Delta}{\sqrt{2}}\right) + e^{\tau/\tau_c}\,{\rm erfc}\!\left(\frac{\sigma_\Delta/\tau_c + \tau/\sigma_\Delta}{\sqrt{2}}\right)\right]. \tag{13}$$

In the weak-jitter limit $\sigma_\Delta \ll \tau_c$, the Gaussian average of $e^{-|\tau-\Delta|/\tau_c}$ over $\Delta \sim \mathcal{N}(0,\sigma_\Delta^2)$ is dominated by $|\Delta| \ll |\tau|$, so one expands around $\Delta = 0$:

$$g^{(1)}_{\rm obs}(\tau) \approx e^{-|\tau|/\tau_c}\left[1 + \frac{\sigma_\Delta^2(\tau)}{2\tau_c^2} + \mathcal{O}(\sigma_\Delta^4)\right]. \tag{14}$$

Here $\sigma_\Delta \equiv \sigma_\Delta(\tau)$ is τ-dependent through $C_\alpha(\tau)$; the expansion is in $\sigma_\Delta(\tau)/\tau_c$ at fixed τ, and its validity requires $\sigma_\Delta(\tau) \ll \tau_c$ at each lag of interest. The correction is positive because $e^{-|\tau|/\tau_c}$ is convex and Jensen's inequality applies. The normalization $g^{(1)}_{\rm obs}(0) = 1$ is preserved exactly.

C. Narrow-Band Limit

High-resolution spectrographs transmit slivers of the spectrum that approach the narrow-band limit. For a quasi-monochromatic source with center frequency ω_0 and angular-frequency bandwidth $\Delta\omega \ll \omega_0$, write $E_{\rm src}(t) = A(t)\, e^{-i\omega_0 t}$ where $A(t)$ varies on the coherence timescale $\tau_c = 1/\Delta\omega = \lambda^2/(2\pi c\,\Delta\lambda)$. For jitter $|\delta t| \ll \tau_c$, the envelope is unaffected, $A(t-\delta t) \approx A(t)$, and the jitter acts purely as a carrier phase modulation:

$$E_{\rm obs}(t) \approx A(t)\, e^{-i\omega_0 t}\, e^{i\omega_0 \delta t(t)}. \tag{15}$$

Since $A(t)$ and $\delta t(t)$ are independent, the first-order coherence factorizes. The jitter contribution is the characteristic function of the Gaussian $\Delta(\tau)$ evaluated at ω_0:

$$\left\langle e^{i\omega_0 \Delta(\tau)} \right\rangle_{\delta t} = e^{-\omega_0^2 \sigma_\Delta^2(\tau)/2}. \tag{16}$$

Substituting $\omega_0 = 2\pi c/\lambda$ and Eq. (7):

$$\boxed{g^{(1)}_{\rm obs}(\tau) = g^{(1)}_{\rm src}(\tau) \cdot \exp\!\left(-\sigma_\phi^2 \left[1 - C_\alpha(\tau; D)\right]\right),} \tag{17}$$

where we define the rms accumulated phase

$$\sigma_\phi \equiv \frac{2\pi\, \sigma_\ell(D)}{\lambda} = \frac{2\pi \sqrt{A_\alpha}\, D^{1-\alpha}\, \ell_{\rm P}^{\alpha}}{\lambda}. \tag{18}$$

We write the foam suppression factor as

$$\mathcal{F}_{\rm foam}(\tau) \equiv e^{-\sigma_\phi^2 [1 - C_\alpha(\tau; D)]}, \tag{19}$$

so $g^{(1)}_{\rm obs}(\tau) = g^{(1)}_{\rm src}(\tau) \cdot \mathcal{F}_{\rm foam}(\tau)$. At $\tau = 0$: $\mathcal{F}_{\rm foam}(0) = 1$. At $\tau \to \infty$: $\mathcal{F}_{\rm foam}(\infty) = e^{-\sigma_\phi^2}$.

D. Connection to Fringe Visibility and Strehl Ratio

1. *Fringe Visibility and the Lieu–Hillman Condition*

The fringe visibility measured in a stellar interferometer at time delay τ is $V(\tau) = |g^{(1)}_{\rm obs}(\tau)|$. From Eq. (17):

$$V(\tau) = V_{\rm src}(\tau) \cdot \mathcal{F}_{\rm foam}(\tau), \tag{20}$$

where $V_{\rm src}(\tau) = |g^{(1)}_{\rm src}(\tau)|$ is the intrinsic source visibility. Lieu & Hillman [5] argue that fringes are erased when the accumulated phase randomizes over $[0, 2\pi)$, i.e. $\sigma_\phi \gtrsim 1$. In our language this is the limit $\mathcal{F}_{\rm foam} \to 0$: the foam suppression factor collapses and $V \to 0$. The Lieu–Hillman condition $\Delta\phi \geq 2\pi$, where $\Delta\phi \simeq 2\pi\sigma_\ell(D)/\lambda$ [1], is precisely $\sigma_\phi \geq 1$. Our result extends this binary criterion to a continuous suppression curve $V(\tau)$ whose shape encodes the temporal decorrelation $C_\alpha(\tau; D)$.

2. *Strehl Ratio*

For spatially uncorrelated Gaussian phase noise of rms σ_ϕ across the telescope aperture, the Strehl ratio (ratio of the observed PSF peak to the diffraction-limited peak) is given by the Maréchal approximation [1]

$$S = e^{-\sigma_\phi^2}, \tag{21}$$

which is valid for $\sigma_\phi \lesssim 1$; for $\sigma_\phi \gg 1$ the PSF peak saturates to zero and the exponential formula is inapplicable in the telescope-optics sense. Within the foam model, however, $\mathcal{F}_{\rm foam}(\tau \to \infty) = e^{-\sigma_\phi^2}$ is an exact result from the Gaussian characteristic function, independent of the Maréchal approximation. The connection between the two is:

$$S = \mathcal{F}_{\rm foam}(\tau \to \infty), \tag{22}$$

which is numerically indistinguishable from zero for $\sigma_\phi \gg 1$ regardless of which interpretation is used. Both the Strehl and the Lieu–Hillman fringe visibility measure the same underlying quantity $\sigma_\phi^2 = (2\pi\sigma_\ell/\lambda)^2$ — the total accumulated phase variance. The difference is observational channel: Strehl and fringes probe the *spatial* phase variation across the aperture, while the temporal $g^{(1)}$ probes the *temporal* phase variation along a single ray as a function of lag.

III. SECOND-ORDER COHERENCE

A. Conditional Siegert Relation

Fix a realization of the jitter process $\delta t(\cdot)$. Conditioned on it, $E_{\rm obs}(t)$ is a deterministic time-warp of a complex Gaussian field, hence is itself complex Gaussian. For any complex Gaussian field the Siegert relation holds exactly, so

$$g^{(2)}_{\rm obs}(\tau \mid \delta t) - 1 = \left| g^{(1)}_{\rm obs}(\tau \mid \delta t) \right|^2 . \tag{23}$$

(A subtlety: conditioned on δt, the field $E_{\rm obs}(t)$ is generally non-stationary, so the conditional two-time correlator depends on the absolute time t as well as the lag τ. The relation (23) is applied as an instantaneous identity at each fixed pair $(t, t+\tau)$; stationarity is restored upon averaging over δt, whose statistics are by assumption stationary.) Given a fixed δt, the differential jitter $\Delta(\tau)$ is a deterministic number, so $g^{(1)}_{\rm obs}(\tau \mid \delta t) = g^{(1)}_{\rm src}(\tau - \Delta(\tau))$. Using the intrinsic Siegert relation $g^{(2)}_{\rm src} - 1 = |g^{(1)}_{\rm src}|^2$:

$$g^{(2)}_{\rm obs}(\tau \mid \delta t) - 1 = \left| g^{(1)}_{\rm src}(\tau - \Delta(\tau)) \right|^2 = g^{(2)}_{\rm src}(\tau - \Delta(\tau)) - 1. \tag{24}$$

B. Unconditional Second-Order Coherence

Averaging over the jitter distribution:

$$\boxed{g^{(2)}_{\rm obs}(\tau) - 1 = \left\langle g^{(2)}_{\rm src}(\tau - \Delta(\tau)) - 1 \right\rangle_{\delta t} = \int_{-\infty}^{\infty} p(\Delta;\tau) \left[g^{(2)}_{\rm src}(\tau - \Delta) - 1 \right] d\Delta.} \tag{25}$$

The structure is identical to Eq. (11): both $g^{(1)}_{\rm obs}$ and $g^{(2)}_{\rm obs}$ are convolutions of their intrinsic counterparts with the *same* Gaussian jitter kernel.

C. Narrow-Band Limit: $g^{(2)}$ is Unchanged to Leading Order

The exact result for $g^{(2)}_{\rm obs}$ is the convolution Eq. (25), which requires no approximation. The narrow-band simplification follows when the jitter $\delta t \ll \tau_c$ does not appreciably shift the amplitude envelope $A(t)$. Since the field is $E_{\rm obs}(t) = A(t)\, e^{-i\omega_0 t}\, e^{i\omega_0 \delta t(t)}$ (with $A(t-\delta t) \approx A(t)$ for slow envelope variations), the intensity is

$$I_{\rm obs}(t) = |E_{\rm obs}(t)|^2 = |A(t - \delta t(t))|^2. \tag{26}$$

In the limit $\delta t/\tau_c \to 0$ we may set $A(t-\delta t) \approx A(t)$, so that to leading order in $\delta t/\tau_c$

$$I_{\rm obs}(t) \approx |A(t)|^2 = I_{\rm src}(t). \tag{27}$$

Consequently, the intensity statistics agree with the intrinsic source to leading order in $\delta t/\tau_c$:

$$g^{(2)}_{\rm obs}(\tau) = g^{(2)}_{\rm src}(\tau) \qquad \text{(leading order in } \delta t/\tau_c), \tag{28}$$

with corrections of order $(\delta t/\tau_c)^2 \sim (\sigma_\ell/c\tau_c)^2$. Note that $\delta t/\tau_c$ and σ_Δ/τ_c are related by $\sigma_\Delta \le \sqrt{2}\,\sigma_\ell/c = \sqrt{2}\,\delta t_{\rm rms}$, so controlling $\delta t/\tau_c$ is the fundamental expansion parameter.

The physical reason is that intensity is blind to the carrier phase $e^{i\omega_0\delta t(t)}$: carrier-phase modulation does not affect $|E_{\rm obs}(t)|^2$. By contrast, $g^{(1)}$ *is* sensitive to this phase because $\omega_0 \delta t \sim \omega_0 \sigma_\ell/c = \sigma_\phi$ can be large even when $\delta t \ll \tau_c$: for $\alpha = 1/2$ at $D = 1$ Gpc and $\lambda = 500$ nm, $\sigma_\phi \approx 280 \gg 1$.

IV. SIEGERT RELATION VIOLATION

A. Motivation: Sensitivity of Temporal Second-Order Correlations

Although the narrow-band jitter leaves $g^{(2)}$ for a point source unchanged at leading order, any observable foam signature in a temporal second-order measurement is controlled by the residual differential-jitter scale $\sigma_\Delta(\tau)$ relative

to the source coherence time τ_c. In the narrow-band limit the relevant sensitivity parameter for the temporal measurement is $\sigma_\Delta(\tau)/\tau_c$, not σ_ϕ directly. Using $c\tau_c = \lambda^2/(2\pi\Delta\lambda)$ (physics convention) and $\omega_0 = 2\pi c/\lambda$, so $\omega_0\tau_c = \lambda/\Delta\lambda$:

$$\frac{\sigma_\Delta(\tau)}{\tau_c} = \frac{\sqrt{2[1-C_\alpha(\tau)]}\,\sigma_\ell(D)}{c\tau_c} = \frac{\sigma_\phi}{\omega_0\tau_c}\cdot\sqrt{2[1-C_\alpha(\tau)]} = \sigma_\phi\cdot\frac{\Delta\lambda}{\lambda}\cdot\sqrt{2[1-C_\alpha(\tau)]}. \tag{29}$$

The suppression factor $\Delta\lambda/\lambda \ll 1$ means that even when $\sigma_\phi \gg 1$ (fringes erased), the temporal HBT ratio can be of order unity or smaller.

TABLE I. Comparison of foam observables at $D = 1$ Gpc, $\lambda = 500$ nm, $\Delta\lambda/\lambda = 10^{-3}$, $A_\alpha = 1$, $C_\alpha \to 0$. The $\alpha = 1/2$ case is observationally excluded [6]; $\alpha = 2/3$ (holographic) is the most sensitive surviving model.

Method	Probe	$\alpha = 1/2$ (excluded)	$\alpha = 2/3$
Strehl / Lieu–Hillman	σ_ϕ	280	2.8×10^{-8}
Temporal $g^{(1)}$, $g^{(2)}$	$\sigma_\phi\sqrt{2[1-C_\alpha]}\cdot\Delta\lambda/\lambda$	0.40	4.0×10^{-11}

B. General Bound

A foundational result for thermal/chaotic light is the Siegert relation, which connects the second-order (intensity–intensity) coherence to the first-order (field–field) coherence: for stationary complex Gaussian fields one has $g^{(2)}(\tau) - 1 = |g^{(1)}(\tau)|^2$. In practice, this means that once $g^{(1)}$ is known, the expected $g^{(2)}$ is fixed, providing a powerful internal consistency check on the underlying statistical assumptions. Systems that do not follow the Siegert relation are therefore of special interest: they signal the presence of nontrivial propagation or measurement physics—e.g., non-Gaussianity, mode mixing, or additional stochastic phase/amplitude processes—that can decouple intensity correlations from field coherence. Such Siegert violations offer a route to probe otherwise hidden dynamics with a self-referential observable, making them particularly valuable in contexts where conventional first-order coherence tests are insensitive or ambiguous.

Define the Siegert discrepancy:

$$\mathcal{S}(\tau) \equiv \left[g^{(2)}_{\rm obs}(\tau) - 1\right] - \left|g^{(1)}_{\rm obs}(\tau)\right|^2. \tag{30}$$

Let $X \equiv g^{(1)}_{\rm src}(\tau - \Delta(\tau))$, a complex random variable. From Eqs. (11) and (25):

$$g^{(1)}_{\rm obs}(\tau) = \langle X\rangle_{\delta t}, \qquad g^{(2)}_{\rm obs}(\tau) - 1 = \langle |X|^2\rangle_{\delta t}. \tag{31}$$

Therefore

$$\mathcal{S}(\tau) = \langle |X|^2\rangle - |\langle X\rangle|^2 \geq 0, \tag{32}$$

by the Cauchy–Schwarz inequality, with equality only when $\Delta(\tau)$ has zero variance (i.e., at $\tau = 0$ or in the absence of foam). The Siegert relation is violated for all $\tau \neq 0$ whenever foam is present. This non-negativity holds for a point source; for extended sources the cross-ray differential jitter can drive $\mathcal{S}(0) < 0$ (see Section V).

C. Explicit Form in the Narrow-Band Limit

In the narrow-band limit, $g^{(2)}_{\rm obs} \approx g^{(2)}_{\rm src}$ to leading order (from Eq. (28)), while $g^{(1)}_{\rm obs}$ is suppressed by $F_{\rm foam}(\tau)$ (Eq. (17)). Using the intrinsic Siegert relation $g^{(2)}_{\rm src}(\tau) - 1 = |g^{(1)}_{\rm src}(\tau)|^2$, to leading order

$$S(\tau) \approx |g^{(1)}_{\rm src}(\tau)|^2\left(1 - e^{-2\sigma_\phi^2[1-C_\alpha(\tau;D)]}\right).$$

D. Properties

The key properties of $\mathcal{S}(\tau)$ are summarized in Table II.

TABLE II. Properties of the Siegert discrepancy $\mathcal{S}(\tau)$ in the narrow-band limit.

Limit	$\mathcal{S}(\tau)$	Reason
$\tau = 0$	0	$\Delta(0) = 0$; both coherences agree
$\sigma_\phi \to 0$ (small phase variance)	$\approx 2\sigma_\phi^2[1 - C_\alpha]\,\lvert g^{(1)}_{\rm src}\rvert^2$	Linear in σ_ϕ^2
$\sigma_\phi \gg 1$ (strong jitter)	$\approx \lvert g^{(1)}_{\rm src}(\tau)\rvert^2$	$g^{(1)}_{\rm obs} \to 0$; full violation
$\tau \to \infty$	$\to 0$	$\lvert g^{(1)}_{\rm src}\rvert \to 0$ faster

E. Physical Interpretation

In the narrow-band limit foam acts as a *phase screen*: it modulates the carrier phase $\omega_0 \delta t(t)$ randomly without touching the intensity envelope. This asymmetry is the root cause of the Siegert violation. The discrepancy $\mathcal{S}(\tau)$ is a direct observable — both $g^{(1)}_{\rm obs}$ and $g^{(2)}_{\rm obs}$ are measurable from the same dataset — and carries a specific distance dependence $\sigma_\phi^2 \propto D^{2(1-\alpha)}$ that distinguishes foam from instrumental effects, which would not scale with source distance. The modulation of $\mathcal{S}(\tau)$ beyond the prefactor $\lvert g^{(1)}_{\rm src}(\tau)\rvert^2$ is controlled solely by α, D, and $C_\alpha(\tau; D)$, independently of the unknown source spectrum.

F. Distinguishing Foam from Conventional Phase Screens

The mathematical structure of Eqs. (17)–(19) is identical to that of any Gaussian random phase screen, including atmospheric turbulence, ionospheric or interstellar scintillation, and intergalactic plasma scattering [4]. Two discriminants separate foam from these conventional sources:

a. Distance scaling. For spacetime foam, $\sigma_\phi^2 \propto D^{2(1-\alpha)}$ (D^1 for $\alpha = 1/2$, $D^{2/3}$ for $\alpha = 2/3$, D^0 for $\alpha = 1$). For a thin intervening screen the phase variance scales as $\sigma_\phi^2 \propto D$ (linear), while for Kolmogorov turbulence distributed along the line of sight the variance accumulates as $D^{11/6}$ [4]. Measuring $\sigma_\phi(D)$ across a range of source redshifts can distinguish these scalings, particularly since the holographic ($\alpha = 2/3$) and conservative ($\alpha = 1$) foam models predict sub-linear or flat growth — qualitatively different from any known astrophysical screen. In practice this program requires knowing the intrinsic source visibility $V_{\rm src}(\tau)$, which is generally unknown and must be modeled or inferred, introducing a degeneracy between foam suppression and intrinsic source properties. Three partial mitigants are available: (i) comparing the same source class at multiple redshifts removes the contribution of intrinsic source variation; (ii) the achromaticity of foam (δt independent of λ) distinguishes it from chromatic intrinsic or propagation effects; and (iii) the Siegert discrepancy $\mathcal{S}(\tau)$ is partially self-calibrating — it measures the difference between $g^{(2)}_{\rm obs}$ and $\lvert g^{(1)}_{\rm obs}\rvert^2$ from the same dataset, so common-mode source properties partially cancel.

b. Frequency dependence. As noted in Section I, geometric-optics foam jitter is achromatic: $\delta t = \Delta\ell/c$ is independent of λ. Plasma-based propagation effects (dispersion measure, scintillation) scale as λ^2, providing a clean observational discriminant when multi-wavelength data are available. Existing upper limits from Fermi and H.E.S.S. time-of-flight measurements [9, 10] already exploit the energy-dependence channel; the coherence-based approach is complementary.

Additional observational constraints on the allowed α range come from imaging coherence measurements: Ragazzoni et al. [8], Perlman et al. [6], and Christiansen et al. [7] together exclude $\alpha \lesssim 2/3$ at optical through X-ray wavelengths for $D \sim 1$–10 Gpc. For surviving models ($\alpha \geq 2/3$), $\sigma_\phi \ll 1$ and all foam coherence observables are proportionally suppressed (see Section III and Table I); the coherence framework then sets upper limits on $\sigma_\phi(D)$ rather than detecting a positive signal.

V. EXTENDED SOURCES: TRANSVERSE FOAM COHERENCE

A. Field Model and Spatiotemporal Covariance

We now generalize to an extended incoherent source whose angular brightness distribution is $W(\boldsymbol{\theta})$ (normalized to unity). The analysis assumes a point (single spatial mode) detector; a finite aperture would introduce Van Cittert–Zernike geometric phases that modify the angular coherence and are left to future work. The detected field is a

weighted superposition over rays labelled by direction offset $\boldsymbol{\theta}$:

$$E_{\rm obs}(t) = \int d\boldsymbol{\theta}\, W^{1/2}(\boldsymbol{\theta})\, E_{\boldsymbol{\theta}}(t - \delta t(t;\boldsymbol{\theta}))\,, \tag{33}$$

where the $E_{\boldsymbol{\theta}}$ are independent, identically distributed complex Gaussian processes (one per sky direction), and each ray carries its own foam jitter $\delta t(t;\boldsymbol{\theta})$. The spatiotemporal covariance of the path-length aggregate is

$$\big\langle \Delta\ell_{\rm tot}(t;\boldsymbol{\theta})\, \Delta\ell_{\rm tot}(t+\tau;\boldsymbol{\theta}')\big\rangle = \sigma_\ell^2(D)\, C_\alpha(\tau;D)\, \Phi\!\left(\frac{b(\boldsymbol{\theta},\boldsymbol{\theta}')}{r_\perp}\right), \tag{34}$$

where $b(\boldsymbol{\theta},\boldsymbol{\theta}') \simeq D|\boldsymbol{\theta}-\boldsymbol{\theta}'|$ is the transverse separation between rays evaluated at the source end, $r_\perp$ is the foam transverse coherence length, and Φ is the transverse coherence function satisfying $\Phi(0) = 1$ and $\Phi(b/r_\perp) \to 0$ for $b \gg r_\perp$. For foam distributed uniformly along the line of sight the ray separation at position z is $z|\boldsymbol{\theta}-\boldsymbol{\theta}'|$, so the full cross-covariance integral involves an average over $0 \le z \le D$; setting $b = D|\boldsymbol{\theta}-\boldsymbol{\theta}'|$ corresponds to a thin-screen approximation in which the correlation is dominated by contributions near the source. Because $r_\perp$ is treated as an effective free parameter, the $O(1)$ geometric factor from the z-integral is absorbed into its definition.

Structurally, $\Phi(b/r_\perp)$ plays the same role in the transverse direction as $C_\alpha(\tau;D)$ plays in the longitudinal (temporal) direction: both are normalized spatial autocorrelation functions of the foam path-length fluctuation field. Just as $\xi_\parallel$ (the longitudinal correlation length) must be inferred from the temporal decorrelation measurements of Sections II and III, the transverse coherence length $r_\perp$ is an independent free parameter of the foam that is constrained empirically by the angular-size dependence of the HBT bunching suppression. Different foam models (random-walk, holographic, etc.) may predict different values of $r_\perp$, but the measurement is model-agnostic in the same way that a turbulence outer-scale measurement constrains L_0 irrespective of the spectral model assumed.

B. First-Order Coherence is Unchanged

The numerator of $g^{(1)}_{\rm obs}(\tau)$ involves cross terms $\langle E^*_{\boldsymbol{\theta}}(\ldots)\, E_{\boldsymbol{\theta}'}(\ldots)\rangle$. Because different source directions are spatially incoherent,

$$\langle E^*_{\boldsymbol{\theta}}(s)\, E_{\boldsymbol{\theta}'}(s')\rangle \propto \delta(\boldsymbol{\theta}-\boldsymbol{\theta}'), \tag{35}$$

so only the diagonal $\boldsymbol{\theta} = \boldsymbol{\theta}'$ contributes. Each diagonal term reproduces the single-ray result of Section II C, and the angular integral over W gives unity. Therefore:

$$g^{(1)}_{\rm obs}(\tau) = g^{(1)}_{\rm src}(\tau) \cdot e^{-\sigma_\phi^2[1-C_\alpha(\tau;D)]}, \tag{36}$$

identical to the point-source result Eq. (17). The first-order coherence probes only single-ray statistics; the transverse foam structure leaves no imprint on $g^{(1)}$.

C. Cross-Ray Differential Jitter

For $g^{(2)}$, the Wick contraction in the Gaussian moment theorem enforces pairings between rays $\boldsymbol{\theta}$ and $\boldsymbol{\theta}'$ that need not be equal. The relevant quantity is the *cross-ray differential jitter*:

$$\Delta_{\boldsymbol{\theta},\boldsymbol{\theta}'}(\tau) \equiv \delta t(t+\tau;\boldsymbol{\theta}) - \delta t(t;\boldsymbol{\theta}'), \tag{37}$$

the time offset between the "second photon" arriving along $\boldsymbol{\theta}$ at $t+\tau$ and the "first photon" along $\boldsymbol{\theta}'$ at t. Computing its variance from Eq. (34):

$$\begin{aligned}\sigma_\Delta^2(\tau;\boldsymbol{\theta},\boldsymbol{\theta}') &= \frac{1}{c^2}\big\langle[\Delta\ell_{\rm tot}(t+\tau;\boldsymbol{\theta}) - \Delta\ell_{\rm tot}(t;\boldsymbol{\theta}')]^2\big\rangle \\ &= \frac{1}{c^2}\left[\sigma_\ell^2 + \sigma_\ell^2 - 2\sigma_\ell^2\, C_\alpha(\tau)\, \Phi\!\left(\frac{b}{r_\perp}\right)\right] \\ &= \frac{2\sigma_\ell^2(D)}{c^2}\left[1 - C_\alpha(\tau;D)\, \Phi\!\left(\frac{b(\boldsymbol{\theta},\boldsymbol{\theta}')}{r_\perp}\right)\right].\end{aligned} \tag{38}$$

Setting $\boldsymbol{\theta} = \boldsymbol{\theta}'$ (and hence $\Phi(0) = 1$) recovers the single-ray result Eq. (7) exactly. For $b \gg r_\perp$, $\Phi \to 0$ and the variance becomes $2\sigma_\ell^2/c^2$ *at all lags, including* $\tau = 0$ — a new effect absent in the point-source case.

D. General Second-Order Coherence

The unnormalized intensity correlator is $\langle I(t)I(t+\tau)\rangle = \langle E^*(t)E(t)E^*(t+\tau)E(t+\tau)\rangle$, where $E \equiv E_{\rm obs}$ is the field of Eq. (33). Expanding the double angular integral and applying the Wick (moment) theorem to the four-point function of the zero-mean complex Gaussian field $E_{\boldsymbol{\theta}}$ produces two pairings (the $W^{1/2}$ weight factors enter the angular integral and are suppressed here for clarity):

$$\langle E^*_{\boldsymbol{\theta}_1} E_{\boldsymbol{\theta}_2} E^*_{\boldsymbol{\theta}_3} E_{\boldsymbol{\theta}_4}\rangle = \langle E^*_{\boldsymbol{\theta}_1} E_{\boldsymbol{\theta}_2}\rangle\langle E^*_{\boldsymbol{\theta}_3} E_{\boldsymbol{\theta}_4}\rangle + \langle E^*_{\boldsymbol{\theta}_1} E_{\boldsymbol{\theta}_4}\rangle\langle E^*_{\boldsymbol{\theta}_3} E_{\boldsymbol{\theta}_2}\rangle. \tag{39}$$

The first pairing contributes $\langle I(t)\rangle\langle I(t+\tau)\rangle$ (the disconnected term, which cancels when forming $g^{(2)}-1$). The second pairing gives the connected contribution. Because different source directions are spatially incoherent, $\langle E^*_{\boldsymbol{\theta}}(s)\, E_{\boldsymbol{\theta}'}(s')\rangle \propto \delta(\boldsymbol{\theta}-\boldsymbol{\theta}')$, so the connected pairing requires $\boldsymbol{\theta}_1 = \boldsymbol{\theta}_4$ and $\boldsymbol{\theta}_3 = \boldsymbol{\theta}_2$. The four-fold angular integral therefore collapses to a double integral over ray pairs $(\boldsymbol{\theta}, \boldsymbol{\theta}')$, each pair contributing its own jitter-convolved $g^{(2)}_{\rm src}$:

$$\boxed{g^{(2)}_{\rm obs}(\tau) - 1 = \iint d\boldsymbol{\theta}\, d\boldsymbol{\theta}'\; W(\boldsymbol{\theta})W(\boldsymbol{\theta}') \int_{-\infty}^{\infty} p_{\boldsymbol{\theta},\boldsymbol{\theta}'}(\Delta;\tau)\left[g^{(2)}_{\rm src}(\tau-\Delta) - 1\right] d\Delta,} \tag{40}$$

where $p_{\boldsymbol{\theta},\boldsymbol{\theta}'}(\Delta;\tau)$ is Gaussian with variance $\sigma^2_\Delta(\tau;\boldsymbol{\theta},\boldsymbol{\theta}')$ from Eq. (38). In the limit $\Phi \equiv 1$ (all rays transversely coherent), every kernel reduces to the single-ray $p_{\rm 1D}$ and Eq. (40) reduces to the point-source result Eq. (25).

E. Effective Transverse Coherence Factor

To reduce Eq. (40) to a tractable form, define the effective transverse coherence factor:

$$\Phi_{\rm eff} \equiv \iint d\boldsymbol{\theta}\, d\boldsymbol{\theta}'\; W(\boldsymbol{\theta})W(\boldsymbol{\theta}')\, \Phi\!\left(\frac{b(\boldsymbol{\theta},\boldsymbol{\theta}')}{r_\perp}\right). \tag{41}$$

This replaces the ray-dependent differential variance with a single effective variance:

$$\sigma^2_{\tau,\rm eff}(\tau) \equiv \frac{2\sigma^2_\ell(D)}{c^2}\Big[1 - C_\alpha(\tau;D)\,\Phi_{\rm eff}\Big], \tag{42}$$

and allows Eq. (40) to be approximated by the simpler 1D convolution of Eq. (25) with σ^2_Δ replaced by $\sigma^2_{\tau,\rm eff}$. The two limits are: $\Phi_{\rm eff} \to 1$ (point-like source or $r_\perp \gg D\,\Delta\theta$) recovers the single-ray result; $\Phi_{\rm eff} \to 0$ (fully incoherent, $r_\perp \ll D\,\Delta\theta$) maximizes the effective variance. The approximation is exact at both limits. At intermediate angular sizes ($r_\perp \sim D\,\Delta\theta$), $\Phi(b/r_\perp)$ varies across ray pairs and the replacement of the full double integral by a single kernel introduces an error of order $\mathrm{Var}_W[\Phi]$, the angular-weighted variance of Φ over the source brightness distribution $W(\boldsymbol{\theta})$. A precise bound on this error requires knowledge of the functional form of Φ and the source profile; for the purposes of this paper the $\Phi_{\rm eff}$ reduction is used as a tractable approximation.

F. Zero-Lag Bunching Peak Suppression

At $\tau = 0$, $C_\alpha(0;D) = 1$, so $\sigma^2_{\tau,\rm eff}(0)$ reduces to

$$\sigma^2_0 \equiv \frac{2\sigma^2_\ell(D)}{c^2}(1 - \Phi_{\rm eff}), \tag{43}$$

which is strictly positive whenever $\Phi_{\rm eff} < 1$. Within the $\Phi_{\rm eff}$ approximation of Section V E, the ray-pair mixture reduces to a single Gaussian kernel of width σ_0. Evaluating the zero-lag convolution using $g^{(2)}_{\rm src}(\tau) - 1 = e^{-2|\tau|/\tau_c}$ (from Siegert applied to $g^{(1)}_{\rm src} = e^{-|\tau|/\tau_c}$): For $\Delta \sim \mathcal{N}(0,\sigma^2_0)$ and $s \equiv \sigma_0/\tau_c$, splitting on $\Delta > 0$ and $\Delta < 0$ and completing the square in each half:

$$\left\langle e^{-2|\Delta|/\tau_c}\right\rangle = \frac{2}{\sqrt{2\pi}\,\sigma_0}\int_0^\infty e^{-2\Delta/\tau_c} e^{-\Delta^2/2\sigma_0^2} d\Delta = e^{2s^2}\,\mathrm{erfc}\!\left(\sqrt{2}\,s\right), \tag{44}$$

which satisfies the limits $s \to 0$: $1 \cdot 1 = 1$ (no suppression) and $s \to \infty$: setting $x = \sqrt{2}s$, the standard asymptotic $e^{x^2}\mathrm{erfc}(x) \sim 1/(x\sqrt{\pi})$ gives $e^{2s^2}\mathrm{erfc}(\sqrt{2}s) \sim 1/(\sqrt{2\pi}\, s) \to 0$ (full suppression, $g^{(2)}_{\mathrm{obs}}(0) \to 1$). so

$$\boxed{g^{(2)}_{\mathrm{obs}}(0) = 1 + e^{2s^2}\,\mathrm{erfc}\Big(\sqrt{2}\,s\Big), \qquad s \equiv \sigma_0/\tau_c,} \tag{45}$$

and the zero-lag bunching peak is suppressed below its ideal value of 2 by

$$\delta g^{(2)} \equiv 2 - g^{(2)}_{\mathrm{obs}}(0) = 1 - e^{2s^2}\,\mathrm{erfc}\Big(\sqrt{2}\,s\Big). \tag{46}$$

In the point-source limit $\Phi_{\mathrm{eff}} \to 1$, $\sigma_0 \to 0$ and $\delta g^{(2)} \to 0$ as expected. Since $\sigma_0^2 \propto (1 - \Phi_{\mathrm{eff}})$ and Φ_{eff} decreases as the source angular diameter increases at fixed D and $r_\perp$, more extended sources experience greater suppression of the bunching peak.

In practice, a real HBT measurement returns $g^{(2)}(0) = 1 + 1/M$ where M is the number of collected spatiotemporal modes (spatial modes times temporal bins within the coherence time), so the foam signal $\delta g^{(2)}$ must exceed the $1/M$ shot-noise floor. For stellar intensity interferometry at optical wavelengths, $M \sim 10^{10}$–10^{12}, giving $\delta g^{(2)}_{\mathrm{floor}} \sim 10^{-10}$–$10^{-12}$. For small s, $\delta g^{(2)} \approx s^2 = (\sigma_0/\tau_c)^2$, so the detectability condition is $\sigma_0/\tau_c \gtrsim M^{-1/2}$; this should be quoted alongside any observational upper bound on σ_0.

The physical origin is distinct from the temporal phase-scrambling effect. Even at $\tau = 0$, photon pairs from *different* sky directions $\boldsymbol{\theta} \neq \boldsymbol{\theta}'$ accumulate independent instantaneous jitters $\delta t(t;\boldsymbol{\theta}) \neq \delta t(t;\boldsymbol{\theta}')$ when $b \gg r_\perp$. Their relative offset $\delta t(t;\boldsymbol{\theta}) - \delta t(t;\boldsymbol{\theta}')$ blurs the $g^{(2)}(0)$ measurement even though $\tau = 0$. This is a purely *spatial* effect with no analogue in the point-source case.

G. Extended-Source Siegert Discrepancy

For the extended source, $g^{(1)}_{\mathrm{obs}}(\tau)$ is unchanged from the point-source result (Eq. 36), while $g^{(2)}_{\mathrm{obs}}(\tau)$ is given by the general mixture Eq. (40). Unlike the point-source case, in the narrow-band limit $g^{(2)}_{\mathrm{obs}}$ is *not* equal to $g^{(2)}_{\mathrm{src}}$: the cross-ray differential jitter modifies the intensity correlation through the envelope timing, even though the carrier phase does not affect individual intensities. The Siegert discrepancy is therefore

$$\mathcal{S}(\tau) = \Big[g^{(2)}_{\mathrm{obs}}(\tau) - 1\Big] - \Big|g^{(1)}_{\mathrm{obs}}(\tau)\Big|^2, \tag{47}$$

where $g^{(2)}_{\mathrm{obs}}(\tau) - 1$ is given by Eq. (40) and $|g^{(1)}_{\mathrm{obs}}(\tau)|^2 = |g^{(1)}_{\mathrm{src}}(\tau)|^2\, e^{-2\sigma_\phi^2[1 - C_\alpha(\tau)]}$. Within the Φ_{eff} approximation, at $\tau = 0$, using $g^{(1)}_{\mathrm{obs}}(0) = 1$ exactly (this follows from the spatial-incoherence argument of Section V and is independent of the Φ_{eff} approximation, which affects only $g^{(2)}_{\mathrm{obs}}$ through the effective kernel width σ_0):

$$\mathcal{S}(0) = \Big[g^{(2)}_{\mathrm{obs}}(0) - 1\Big] - 1 = e^{2s^2}\,\mathrm{erfc}\Big(\sqrt{2}\,s\Big) - 1 = -\,\delta g^{(2)} < 0, \tag{48}$$

so within this approximation the Siegert discrepancy at zero lag equals minus the bunching-peak suppression, and is negative for $s > 0$. For a point source ($\Phi_{\mathrm{eff}} = 1$, $s = 0$), $\mathcal{S}(0) = 0$, recovering the point-source result.

The negative value of $\mathcal{S}(0)$ does not violate any fundamental inequality. The bound $\mathcal{S}(\tau) \geq 0$ was established in Section IV for a point source, where the conditional Siegert relation and Cauchy–Schwarz apply to a single complex Gaussian mode conditioned on one jitter process. For an extended source the unconditional field is an *angular mixture* over many independent modes and jitter processes; such a mixture is not itself a Gaussian state in the mode-counting sense, and therefore need not satisfy $\mathcal{S} \geq 0$. Physically, $g^{(2)}_{\mathrm{obs}}(0)$ is suppressed below 2 by the cross-ray spatial jitter (an envelope-level effect), while $|g^{(1)}_{\mathrm{obs}}(0)|^2 = 1$ exactly, so these two observables are controlled by different physical mechanisms and their difference can be negative.

H. Two Distinct Signatures

The two observable signatures of foam for extended sources are summarized in Table III.

The two signatures are complementary. For a point-like source, the Siegert discrepancy $\mathcal{S}(\tau)$ is controlled by σ_ϕ^2 and $C_\alpha(\tau; D)$ alone, probing α and D with no sensitivity to $r_\perp$. For an extended source, $\mathcal{S}(\tau)$ acquires additional

TABLE III. Two foam signatures for an extended source in the narrow-band limit.

Signature	Origin	Non-zero at $\tau = 0$?	Depends on $r_\perp$?
Siegert discrepancy $\mathcal{S}(\tau)$	Temporal + cross-ray jitter	Yes (via $\delta g^{(2)}$; < 0)	Yes — via $\Phi_{\rm eff}$
Bunching suppression $\delta g^{(2)}$	Transverse spatial decorrelation	Yes	Yes — direct probe of $r_\perp$

dependence on $r_\perp$ through the cross-ray differential jitter: at $\tau = 0$ this is captured (within the $\Phi_{\rm eff}$ reduction) by $\mathcal{S}(0) = -\delta g^{(2)}$, which depends on $\Phi_{\rm eff}$ and hence on $r_\perp$. The bunching suppression $\delta g^{(2)}$ is controlled by $\sigma_0^2 \propto \sigma_\ell^2(D)(1 - \Phi_{\rm eff})$, providing a direct handle on $r_\perp$ through the dependence of $\Phi_{\rm eff}$ on source angular size. Joint measurements of point-like and extended sources at the same redshift therefore allow $\{\alpha, \sigma_\ell(D), r_\perp\}$ to be constrained simultaneously and largely independently of the unknown intrinsic source spectrum.

VI. CONCLUSIONS

For $\alpha = 1/2$ at $D = 1$ Gpc with $A_\alpha = 1$: $\sigma_\ell = \sqrt{\ell_{\rm P}\, D} \approx 2.2 \times 10^{-5}$ m, giving $\sigma_\phi \approx 280$ (fringes completely erased at $\sigma_\phi \gg 1$), while $c\tau_c = \lambda^2/(2\pi\Delta\lambda) \approx 80\,\mu$m and $\sigma_\Delta/\tau_c \approx \sqrt{2}\,\sigma_\ell/(c\tau_c) \approx 0.40$. The narrow-band expansion parameter is $\delta t_{\rm rms}/\tau_c = \sigma_\ell/(c\tau_c) \approx 0.28$, so corrections to Eq. (28) are at the $\sim 8\%$ level; the exact convolution (25) should be used for precision work. Critically, this case ($\alpha = 1/2$) is already excluded by fringe-visibility measurements [6–8]. For the holographic model $\alpha = 2/3$ with the same parameters, $\sigma_\phi \sim 10^{-8}$, so all foam coherence observables are negligibly small with current instrumentation. The theoretical framework is therefore most useful as a basis for future constraints at shorter wavelengths (X-ray or near-UV), longer baselines, or next-generation intensity interferometers, where the sensitivity scales as λ^{-1} and $D^{1-\alpha}$. To make this concrete: for $\alpha = 2/3$ at $D = 10$ Gpc and soft X-ray wavelength $\lambda = 1$ Å, one finds $\sigma_\ell \approx 4.3 \times 10^{-15}$ m, $\sigma_\phi \approx 2.7 \times 10^{-4}$, and $\sigma_\Delta/\tau_c \approx 3.8 \times 10^{-7}$ (for $\Delta\lambda/\lambda = 10^{-3}$) — still approximately seven orders of magnitude below a unity HBT signal. The framework therefore serves primarily as a theoretical tool and as an upper-limit setter for future precision measurements, rather than as a near-term positive-detection strategy.

A deeper question left open by the present framework is whether spacetime foam should be treated as a static medium or whether its statistical properties evolve on cosmological timescales. Throughout this paper we have treated $\sigma_\ell^2(D)$ as a fixed parameter, an assumption justified by the negligible fractional change $H_0 T \sim 10^{-10}$ in the propagation distance over any realistic observation duration. Over cosmological lookback times, however, the foam may track the expansion of the universe in a non-trivial way. If the Ng–Perlman exponent α or the coefficient A_α itself carries a redshift dependence, then observations of sources at different redshifts would no longer trace a simple power-law $\sigma_\phi(D)$, and the distance-scaling discriminant developed in Section IV.E would need to be generalized. More speculatively, if the foam undergoes a phase transition or qualitative change in its structure at some epoch — analogous to other symmetry-breaking transitions in the early universe — this could in principle be imprinted on the coherence of light from sources bracketing that epoch, providing a novel cosmological probe. We flag these possibilities as motivation for developing a fully non-stationary, redshift-resolved treatment of foam coherence in future work.

ACKNOWLEDGMENTS

This work was supported in part by the Director, Office of Science, Office of High Energy Physics of the U.S. Department of Energy under Contract No. DE-AC02-05CH11231.

[1] Y. J. Ng and E. S. Perlman, Universe **8**, 382 (2022).
[2] Y. J. Ng, Modern Physics Letters A **18**, 1073 (2003).
[3] Y. J. Ng, Physical Review Letters **86**, 2946 (2001).
[4] S. Carlip, Rep. Prog. Phys. **86**, 066001 (2023), arXiv:2209.14282 [gr-qc].
[5] R. Lieu and L. W. Hillman, The Astrophysical Journal **585**, L77 (2003).
[6] E. S. Perlman, S. A. Rappaport, W. A. Christiansen, Y. J. Ng, J. DeVore, and D. Pooley, The Astrophysical Journal **805**, 10 (2015).
[7] W. A. Christiansen, Y. J. Ng, D. J. E. Floyd, and E. S. Perlman, Physical Review D **83**, 084003 (2011).
[8] R. Ragazzoni, M. Turatto, and W. Gaessler, ApJ **587**, L1 (2003), arXiv:astro-ph/0303043 [astro-ph].

[9] A. A. Abdo, M. Ackermann, M. Ajello, *et al.* (Fermi GBM/LAT Collaboration), Nature **462**, 331 (2009), arXiv:0908.1832 [astro-ph].
[10] F. Aharonian *et al.* (H.E.S.S. Collaboration), Phys. Rev. Lett. **101**, 170402 (2008), arXiv:0810.3475 [astro-ph].